

Forecast of the loss of clients in the Bank

Content

1. [Project Description](#)
2. [Data Preparation](#)
3. [Models training](#)
4. [Combat to imbalance](#)
5. [Models testing](#)
6. [General conclusion](#)

Project Description

Company "Beta bank" has faced the loss of the clients. The clients are started to terminate the contracts with the bank every month. The quantity of clients who terminated the contact with the bank is not sufficient compare to the total quantity of clients. Marketing department of bank calculated that keep the current clients is cheaper than to pull the new one to come.

It's required to make a prediction of the loss of clients in nearest time. The historical data provided and has the information on clients behavior and contract termination with bank.

It's required to train the model and get the **F1** score as high as possible. The value of **F1** equal to 0.59 or higher should be achieved.

Additionally it's required to check *AUC-ROC* scores of models and compare it.

Data original source: <https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling>

Data preparation

Libraries import

```
In [1]: import pandas as pd
        from sklearn.linear_model import LogisticRegression
        from sklearn.ensemble import RandomForestClassifier
        from sklearn.tree import DecisionTreeClassifier
        from sklearn.ensemble import RandomForestRegressor
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import confusion_matrix
        from sklearn.metrics import recall_score
        from sklearn.metrics import precision_score
        from sklearn.metrics import f1_score
        import matplotlib.pyplot as plt
        from sklearn.metrics import roc_curve
        from sklearn.metrics import roc_auc_score
        from sklearn.metrics import r2_score
        from sklearn.metrics import mean_absolute_error
        from sklearn.metrics import accuracy_score
        from sklearn.utils import shuffle

        from sklearn.model_selection import RandomizedSearchCV
        from sklearn.model_selection import GridSearchCV
        import numpy as np
```


```
In [2]: import warnings
        warnings.simplefilter(action='ignore', category=FutureWarning)
        warnings.simplefilter(action='ignore', category=DeprecationWarning)
        warnings.simplefilter(action='ignore', category=RuntimeWarning)
```

Data loading and overview

```
In [3]: data = pd.read_csv('Churn.csv')
        data.head(15)
```

Out[3]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	Estima
0	1	15634602	Hargrave	619	France	Female	42	2.0	0.00	1	1	1	
1	2	15647311	Hill	608	Spain	Female	41	1.0	83807.86	1	0	1	
2	3	15619304	Onio	502	France	Female	42	8.0	159660.80	3	1	0	
3	4	15701354	Boni	699	France	Female	39	1.0	0.00	2	0	0	
4	5	15737888	Mitchell	850	Spain	Female	43	2.0	125510.82	1	1	1	
5	6	15574012	Chu	645	Spain	Male	44	8.0	113755.78	2	1	0	
6	7	15592531	Bartlett	822	France	Male	50	7.0	0.00	2	1	1	
7	8	15656148	Obinna	376	Germany	Female	29	4.0	115046.74	4	1	0	
8	9	15792365	He	501	France	Male	44	4.0	142051.07	2	0	1	
9	10	15592389	H?	684	France	Male	27	2.0	134603.88	1	1	1	
10	11	15767821	Bearce	528	France	Male	31	6.0	102016.72	2	0	0	
11	12	15737173	Andrews	497	Spain	Male	24	3.0	0.00	2	1	0	
12	13	15632264	Kay	476	France	Female	34	10.0	0.00	2	1	0	
13	14	15691483	Chin	549	France	Female	25	5.0	0.00	2	0	0	
14	15	15600882	Scott	635	Spain	Female	35	7.0	0.00	2	1	1	



In [4]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   RowNumber       10000 non-null  int64
1   CustomerId      10000 non-null  int64
2   Surname         10000 non-null  object
3   CreditScore     10000 non-null  int64
4   Geography       10000 non-null  object
5   Gender         10000 non-null  object
6   Age            10000 non-null  int64
7   Tenure         9091 non-null   float64
8   Balance        10000 non-null  float64
9   NumOfProducts  10000 non-null  int64
10  HasCrCard       10000 non-null  int64
11  IsActiveMember  10000 non-null  int64
12  EstimatedSalary 10000 non-null  float64
13  Exited         10000 non-null  int64
dtypes: float64(3), int64(8), object(3)
memory usage: 1.1+ MB
```

```
In [5]: data.sort_values(by='Tenure')
```

Out[5]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
6546	6547	15633840	Henderson	781	France	Male	20	0.0	125023.10	2	1	1
2622	2623	15787026	Onwuatuegwu	627	Germany	Male	27	0.0	185267.45	2	1	1
886	887	15591711	Sleeman	739	Spain	Male	38	0.0	128366.44	1	1	0
4532	4533	15739194	Manfrin	548	Spain	Male	38	0.0	178056.54	2	1	0
8213	8214	15748352	Endrizzi	598	Spain	Male	34	0.0	104488.17	1	0	1
...
9944	9945	15703923	Cameron	744	Germany	Male	41	NaN	190409.34	2	1	1
9956	9957	15707861	Nucci	520	France	Female	46	NaN	85216.61	1	1	0
9964	9965	15642785	Douglas	479	France	Male	34	NaN	117593.48	2	0	0
9985	9986	15586914	Nepean	659	France	Male	36	NaN	123841.49	2	1	0
9999	10000	15628319	Walker	792	France	Female	28	NaN	130142.79	1	1	0

10000 rows × 14 columns



In [6]: `data.query('Tenure == "NaN"')['RowNumber'].count()`

Out[6]: 0

In [7]: `data['Tenure'] = data['Tenure'].fillna(0)
data.sort_values(by='Tenure')`

Out[7]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
9999	10000	15628319	Walker	792	France	Female	28	0.0	130142.79	1	1	0	
6220	6221	15716926	Macleod	807	France	Male	33	0.0	101952.97	2	1	0	
6221	6222	15603554	Berkeley	513	France	Female	45	0.0	164649.52	3	1	0	
6223	6224	15679429	Bell	694	France	Male	32	0.0	91956.49	1	1	1	
6225	6226	15742172	Williamson	598	Germany	Male	32	0.0	123938.60	2	1	0	
...
7579	7580	15649101	Reeves	601	France	Male	40	10.0	127847.86	1	0	0	
7565	7566	15623369	Clifton	708	France	Male	52	10.0	105355.81	1	1	0	
1619	1620	15770309	McDonald	656	France	Male	18	10.0	151762.74	1	0	1	
3135	3136	15753874	Kent	694	France	Male	37	10.0	143835.47	1	0	1	
4863	4864	15640491	Raff	464	France	Female	33	10.0	147493.70	2	1	0	

10000 rows × 14 columns



In [8]:

```
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   RowNumber       10000 non-null  int64
1   CustomerId      10000 non-null  int64
2   Surname         10000 non-null  object
3   CreditScore     10000 non-null  int64
4   Geography       10000 non-null  object
5   Gender          10000 non-null  object
6   Age            10000 non-null  int64
7   Tenure          10000 non-null  float64
8   Balance         10000 non-null  float64
9   NumOfProducts  10000 non-null  int64
10  HasCrCard       10000 non-null  int64
11  IsActiveMember  10000 non-null  int64
12  EstimatedSalary 10000 non-null  float64
13  Exited          10000 non-null  int64
dtypes: float64(3), int64(8), object(3)
memory usage: 1.1+ MB

```

```

In [9]: # deletion of useless columns
data = data.drop(columns = ['RowNumber', 'CustomerId', 'Surname'])

```

```

In [10]: data.head()

```

```

Out[10]:
   CreditScore  Geography  Gender  Age  Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember  EstimatedSalary  Exited
0          619     France  Female   42    2.0     0.00             1           1              1         101348.88         1
1          608      Spain  Female   41    1.0  83807.86             1           0              1         112542.58         0
2          502     France  Female   42    8.0  159660.80             3           1              0         113931.57         1
3          699     France  Female   39    1.0     0.00             2           0              0          93826.63         0
4          850      Spain  Female   43    2.0  125510.82             1           1              1          79084.10         0

```

Conclusion

1) Data successfully imported;

- 2) Useles columns were deleted;
- 3) nulls filled with zeros;
- 4) prepared dataset has 11 columns and 10000 rows.

Models training

Encoding of categorical data

```
In [11]: data_s = data.copy()
```

```
In [12]: data_s = pd.get_dummies(data_s,drop_first=True)
```

```
In [13]: data_s.head(10)
```

```
Out[13]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited	Geography_Germany	Geography_Spain
0	619	42	2.0	0.00	1	1	1	101348.88	1	0	0
1	608	41	1.0	83807.86	1	0	1	112542.58	0	0	1
2	502	42	8.0	159660.80	3	1	0	113931.57	1	0	0
3	699	39	1.0	0.00	2	0	0	93826.63	0	0	0
4	850	43	2.0	125510.82	1	1	1	79084.10	0	0	1
5	645	44	8.0	113755.78	2	1	0	149756.71	1	0	1
6	822	50	7.0	0.00	2	1	1	10062.80	0	0	0
7	376	29	4.0	115046.74	4	1	0	119346.88	1	1	0
8	501	44	4.0	142051.07	2	0	1	74940.50	0	0	0
9	684	27	2.0	134603.88	1	1	1	71725.73	0	0	0


```
In [14]: features = data_s.drop(columns='Exited')
        target = data_s['Exited']
```

Dataset splitting on samples

```
In [15]: features_train, features_valid_temp, target_train, target_valid_temp = train_test_split(
        features, target, test_size=0.4, random_state=12345)
```

```
In [16]: features_valid, features_test, target_valid, target_test = train_test_split(features_valid_temp, target_valid_temp, test_size=0.5, ra
```

Random forest Model training

```
In [17]: rf_model = RandomForestClassifier(random_state=12345)
        rf_model.fit(features_train, target_train)
        rf_predictions = rf_model.predict(features_valid)
```

```
In [18]: confusion_matrix(target_valid, rf_predictions)
```

```
Out[18]: array([[1521,   61],
        [ 222,  196]], dtype=int64)
```

```
In [19]: rf_f1_score = f1_score(target_valid, rf_predictions)
        rf_f1_score
```

```
Out[19]: 0.5807407407407408
```

```
In [20]: recall_score(target_valid, rf_predictions)
```

```
Out[20]: 0.4688995215311005
```

```
In [21]: precision_score(target_valid, rf_predictions)
```

```
Out[21]: 0.7626459143968871
```


```
In [22]: accuracy_score(target_valid, rf_predictions)
```

```
Out[22]: 0.8585
```

Plotting of random forest model auc roc curve

```
In [23]: rf_probabilities_valid = rf_model.predict_proba(features_valid)
precision, recall, thresholds = roc_curve(target_valid, rf_probabilities_valid[:, 1])

plt.figure(figsize=(6, 6))
plt.step(precision, recall, where='post')
plt.plot([0.0, 1.0], [0.0, 1.0])
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.title('AUC ROC curve')
plt.show()
```



```
In [24]: roc_auc_score(target_valid, rf_probabilities_valid[:, 1])
```

```
Out[24]: 0.8424114893025
```

Conclusions:

Random Forest model was successfully trained and has the following scores: 1) F1 - 0.53

2) AUC ROC - 0.81

3) recall_score - 0.418

4) precision_score - 0.73

5) accuracy_score - 0.84

Combat to imbalance

Logistic Regression model training with balanced weight

```
In [25]: lr_model = LogisticRegression(random_state=12345, solver='liblinear', class_weight='balanced')  
lr_model.fit(features_train, target_train)  
lr_predictions = lr_model.predict(features_valid)
```

```
In [26]: lr_f1_score = f1_score(target_valid, lr_predictions)  
lr_f1_score
```

```
Out[26]: 0.4514056224899599
```

```
In [27]: recall_score(target_valid, lr_predictions)
```

```
Out[27]: 0.6722488038277512
```

```
In [28]: precision_score(target_valid, lr_predictions)
```

```
Out[28]: 0.33978234582829503
```

```
In [29]: accuracy_score(target_valid, lr_predictions)
```

Out[29]: 0.6585


```
In [30]: confusion_matrix(target_valid,lr_predictions)
```

```
Out[30]: array([[1036,  546],
               [ 137,  281]], dtype=int64)
```

Plotting of logistic regression model auc roc curve

```
In [31]: lr_probabilities_valid = lr_model.predict_proba(features_valid)
precision, recall, thresholds = roc_curve(target_valid, lr_probabilities_valid[:, 1])

plt.figure(figsize=(6, 6))
plt.step(precision,recall, where='post')
plt.plot([0.0,1.0],[0.0,1.0])
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.title('Curve AUC ROC')
plt.show()
```



```
In [32]: roc_auc_score(target_valid,lr_probabilities_valid[:, 1])
```

```
Out[32]: 0.7193909955903435
```

Conclusion

Logistic regression model with balanced class weight has the following scores:

- 1) F1 - 0.48
- 2) AUC ROC - 0.75
- 3) recall_score - 0.37
- 4) precision_score - 0.73

5) accuracy_score - 0.69

Upsampling

```
In [33]: features_zeros = features_train[target_train==0]
features_ones = features_train[target_train==1]
target_zeros = target_train[target_train==0]
target_ones = target_train[target_train==1]
n= 15
```

```
In [34]: features_upsampled = pd.concat([features_zeros]+[features_ones]*n)
target_upsampled = pd.concat([target_zeros]+[target_ones]*n)
```

```
In [35]: features_upsampled,target_upsampled = shuffle(features_upsampled,target_upsampled,random_state=12345)
```

Random forest model training on upsampled data

```
In [36]: rf_up_model = RandomForestClassifier(random_state=12345)
rf_up_model.fit(features_upsampled,target_upsampled)
rf_up_predictions= rf_up_model.predict(features_valid)
rf_up_f1_score = f1_score(target_valid, rf_up_predictions)
rf_up_f1_score
```

```
Out[36]: 0.5932885906040269
```

```
In [37]: confusion_matrix(target_valid,rf_up_predictions)
```

```
Out[37]: array([[1476, 106],
               [ 197, 221]], dtype=int64)
```

```
In [38]: recall_score(target_valid, rf_up_predictions)
```

```
Out[38]: 0.5287081339712919
```


```
In [39]: precision_score(target_valid, rf_up_predictions)
```

```
Out[39]: 0.6758409785932722
```

Plotting of upsampled rf model auc roc curve

```
In [40]: rf_up_probabilities_valid = rf_up_model.predict_proba(features_valid)
precision, recall, thresholds = roc_curve(target_valid, rf_up_probabilities_valid[:, 1])

plt.figure(figsize=(6, 6))
plt.step(precision, recall, where='post')
plt.plot([0.0, 1.0], [0.0, 1.0])
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.title('Curve AUC ROC')
plt.show()
```



```
In [41]: roc_auc_score(target_valid, rf_up_probabilities_valid[:, 1])
```

```
Out[41]: 0.8381833001651353
```

```
In [42]: accuracy_score(target_valid, rf_up_predictions)
```

```
Out[42]: 0.8485
```

Conclusion

Random forest model on upsampled data has the following scores:

- 1) F1 - 0.57
- 2) AUC ROC - 0.808
- 3) recall_score - 0.50
- 4) precision_score - 0.65
- 5) accuracy_score - 0.84

Downsampling

```
In [43]: features_zeros = features_train[target_train==0]
features_ones = features_train[target_train==1]
target_zeros = target_train[target_train==0]
target_ones = target_train[target_train==1]
m= 0.02
```

```
In [44]: features_downsampled = pd.concat(
        [features_zeros.sample(frac=m, random_state=12345)] + [features_ones])
target_downsampled = pd.concat(
        [target_zeros.sample(frac=m, random_state=12345)] + [target_ones])
```

```
In [45]: features_downsampled, target_downsampled = shuffle(features_upsampled, target_upsampled, random_state=12345)
```

```
In [46]: rf_dn_model = RandomForestClassifier(random_state=12345)
rf_dn_model.fit(features_downsampled, target_downsampled)
rf_dn_predictions= rf_dn_model.predict(features_valid)
rf_dn_f1_score = f1_score(target_valid, rf_dn_predictions)
rf_dn_f1_score
```


Out[46]: 0.5850340136054422

```
In [47]: confusion_matrix(target_valid, rf_dn_predictions)
```

```
Out[47]: array([[1480,  102],
               [ 203,  215]], dtype=int64)
```

```
In [48]: rf_dn_probabilities_valid = rf_dn_model.predict_proba(features_valid)
precision_score(target_valid, rf_dn_predictions)
```

Out[48]: 0.6782334384858044

```
In [49]: recall_score(target_valid, rf_dn_predictions)
```

Out[49]: 0.5143540669856459

```
In [50]: roc_auc_score(target_valid, rf_dn_probabilities_valid[:, 1])
```

Out[50]: 0.837675947713209


```
In [51]: accuracy_score(target_valid, rf_dn_predictions)
```

Out[51]: 0.8475

Plotting of downsampled rf model auc roc curve

```
In [52]: probabilities_valid = rf_dn_model.predict_proba(features_valid)
precision, recall, thresholds = roc_curve(target_valid, probabilities_valid[:, 1])

plt.figure(figsize=(6, 6))
plt.step(precision, recall, where='post')
plt.plot([0.0, 1.0], [0.0, 1.0])
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.title('Curve AUC ROC')
plt.show()
```



Conclusion

Random forest model with downsampling has the following scores:

- 1) F1 - 0.55
- 2) AUC ROC - 0.81
- 3) recall_score - 0.47
- 4) precision_score - 0.66
- 5) accuracy_score - 0.84

Hyperparameters tuning

In [53]: %%time

```
models = [RandomForestClassifier,DecisionTreeClassifier]

f1_best=0
for c in models:
    for d in range(2):
        if d == 0:
            e = features_downsampled
            f = target_downsampled
            g = 'downsampled'
        elif d == 1:
            e = features_upsampled
            f = target_upsampled
            g = 'upsampled'
        else:
            e = features_train
            f = target_train
            g = 'train'
        for depth in range (10,210,50):
            if c == RandomForestClassifier:
                for estimators in range (10,210,50):
                    tuned_model = c(n_estimators = estimators , max_depth = depth, random_state=12345,
                                    class_weight ='balanced')
                    tuned_model.fit(e,f)
                    tuned_predictions= tuned_model.predict(features_valid)
                    F1_temp = f1_score(target_valid, tuned_predictions)
                    if f1_best < F1_temp:
                        f1_best = F1_temp
                        best_model = c
                        best_features = g
                        best_depth = depth
                        best_estimators = estimators
            elif c == DecisionTreeClassifier:
                tuned_model = c(max_depth = depth, random_state=12345,class_weight ='balanced')
                tuned_model.fit(e,f)
                tuned_predictions= tuned_model.predict(features_valid)
                F1_temp = f1_score(target_valid, tuned_predictions)
                if f1_best < F1_temp:
                    f1_best = F1_temp
                    best_model = c
                    best_features = g
                    best_depth = depth
```

```
best_estimators = estimators
print(f1_best, '\n', best_model, '\n', best_features, '\n', best_depth, '\n', best_estimators)
```

```
0.6293103448275862
```

```
<class 'sklearn.ensemble._forest.RandomForestClassifier'>
```

```
upsampled
```

```
10
```

```
60
```

```
CPU times: total: 41 s
```

```
Wall time: 41.1 s
```

```
In [54]: tuned_model = best_model(n_estimators = best_estimators, max_depth =best_depth, random_state=12345, class_weight = 'balanced')
tuned_model.fit(e, f)
tuned_predictions = tuned_model.predict(features_valid)
```

```
In [55]: tuned_f1 = f1_score(target_valid, tuned_predictions)
tuned_f1
```

```
Out[55]: 0.6293103448275862
```

```
In [56]: accuracy_score(target_valid, tuned_predictions)
```

```
Out[56]: 0.828
```

```
In [57]: recall_score(target_valid, tuned_predictions)
```

```
Out[57]: 0.6985645933014354
```

```
In [58]: tuned_probabilities_valid = tuned_model.predict_proba(features_valid)
roc_auc_score(target_valid,tuned_probabilities_valid[:, 1])
```

```
Out[58]: 0.8502728059085768
```


```
In [59]: precision_score(target_valid, tuned_predictions)
```

```
Out[59]: 0.5725490196078431
```

Plotting of AUC ROC curve of tuned model

```
In [60]: tuned_probabilities_valid = tuned_model.predict_proba(features_valid)
precision, recall, thresholds = roc_curve(target_valid, tuned_probabilities_valid[:, 1])
```

```
plt.figure(figsize=(6, 6))
plt.step(precision, recall, where='post')
plt.plot([0.0, 1.0], [0.0, 1.0])
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.title('Curve AUC ROC')
plt.show()
```



Conclusions

Using hyperparameters tuning the model has the following scores:

1) F1 ~ 0.629

2) AUC ROC ~ 0.85

3) recall_score ~ 0.44

4) precision_score ~ 0.57

5) accuracy_score ~ 0.828

Selection of best model

```
In [61]: models_df = pd.DataFrame({'model_name': ['rf_model', 'lr_model', 'rf_up_model', 'rf_dn_model', 'tuned_model'],  
                                'model': [rf_model, lr_model, rf_up_model, rf_dn_model, tuned_model],  
                                'F1': [rf_f1_score, lr_f1_score, rf_up_f1_score, rf_dn_f1_score, tuned_f1] })
```

```
In [62]: models_df = models_df.sort_values(by = 'F1', ascending = 0).reset_index(drop = True)  
models_df
```

```
Out[62]:
```

	model_name	model	F1
0	tuned_model	(DecisionTreeClassifier(max_depth=10, max_feat...	0.629310
1	rf_up_model	(DecisionTreeClassifier(max_features='sqrt', r...	0.593289
2	rf_dn_model	(DecisionTreeClassifier(max_features='sqrt', r...	0.585034
3	rf_model	(DecisionTreeClassifier(max_features='sqrt', r...	0.580741
4	lr_model	LogisticRegression(class_weight='balanced', ra...	0.451406

```
In [63]: best_model = models_df['model'][0]  
best_model
```

```
Out[63]:
```

▼ RandomForestClassifier

RandomForestClassifier(class_weight='balanced', max_depth=10, n_estimators=60,
random_state=12345)

Conclusion

1) Five models were trained and compared:

- randomforest model;

- linear regression with class balance;
- random forest with class upsampling;
- random forest with class downsampling;
- random forest with hyperparameters tuning.

2) Best F1 score was achieved using tuned model. The model was selected for testing.

Model testing

```
In [64]: best_model.fit(features_upsampled, target_upsampled)
test_predictions = best_model.predict(features_test)
print( 'Test_f1:', f1_score(target_test, test_predictions))
```

```
Test_f1: 0.5982905982905982
```

```
In [65]: accuracy_score(target_test, test_predictions)
```

```
Out[65]: 0.812
```

```
In [66]: recall_score(target_test, test_predictions)
```

```
Out[66]: 0.6619385342789598
```

```
In [67]: test_probabilities = best_model.predict_proba(features_test)
roc_auc_score(target_test, test_probabilities[:, 1])
```

```
Out[67]: 0.8544157968192292
```


```
In [68]: precision_score(target_test, test_predictions)
```

```
Out[68]: 0.5458089668615984
```

```
In [69]: test_probabilities = best_model.predict_proba(features_test)
precision, recall, thresholds = roc_curve(target_test, test_probabilities[:, 1])

plt.figure(figsize=(6, 6))
```

```
plt.step(precision,recall, where='post')
plt.plot([0.0,1.0],[0.0,1.0])
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.title('Curve AUC ROC')
plt.show()
```



Conclusion

Testing of the model with highest F1 score on valid sample has the following scores:

- 1) F1 ~ 0.598
- 2) AUC ROC ~ 0.85
- 3) recall_score ~ 0.67

4) precision_score ~ 0.54

5) accuracy_score ~ 0.81

General conclusion

1) Data was successfully imported, prepared and encoded.

2) Dataset was splitted on train, valid and test samples.

3) Five models were train, four of them were using the different tools for combat to imbalance of classes:

- balancing of weight of classes;
- upsampling;
- down sampling;

4) The best model was selected after comparison. The model has tuned hyperparameters and balanced class weight.

5) The main goal was achieved - the F1 score is above 0.59.