

Prediction of the star temperature

Content

1. [Project Description](#)
2. [Data import](#)
3. [Data preparation and analysis](#)
4. [Data Encoding](#)
5. [Neural networks building](#)
6. [Neural network tuning](#)
7. [General Conclusion](#)

Project Description

Astronomy company would like to define the temperature on the surface of the newly spotted stars. For that purpose it's required to predict the star temperature using machine learning and based on the previous data from the database of the stars.

Main tasks are:

- import and prepare data;
- perform the exploratory analysis;
- rescale the data, and split it on train and test samples;
- build neural networks;
- declare the function for training and prediction;
- plot the comparison chart of actual and predicted star temperature;
- perform the network tuning (using dropout, optimisers, quantity of layers, activation functions);
- compare the models and select the best;

- RMSE score on the best network shall be 4500 or lower.

Data import

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
from sklearn.preprocessing import OneHotEncoder, StandardScaler, MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
```

```
In [10]: df_stars = pd.read_csv('data_stars.csv', index_col = [0])
```

Data preparation and analysis

```
In [11]: # display 5 first rows
df_stars.head()
```

```
Out[11]:
```

	Temperature (K)	Luminosity(L/Lo)	Radius(R/Ro)	Absolute magnitude(Mv)	Star type	Star color
0	3068	0.002400	0.1700	16.12	0	red
1	3042	0.000500	0.1542	16.60	0	red
2	2600	0.000300	0.1020	18.70	0	red
3	2800	0.000200	0.1600	16.65	0	red
4	1939	0.000138	0.1030	20.06	0	red

```
In [13]: # display info
df_stars.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 228 entries, 0 to 239
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Temperature (K)        228 non-null   int64
1   Luminosity(L/Lo)       228 non-null   float64
2   Radius(R/Ro)           228 non-null   float64
3   Absolute magnitude(Mv)  228 non-null   float64
4   Star type              228 non-null   int64
5   Star color             228 non-null   object
dtypes: float64(3), int64(2), object(1)
memory usage: 12.5+ KB

```

```

In [14]: # display the data of numeric columns
df_stars.describe()

```

```

Out[14]:

```

	Temperature (K)	Luminosity(L/Lo)	Radius(R/Ro)	Absolute magnitude(Mv)	Star type
count	228.000000	228.000000	228.000000	228.000000	228.000000
mean	10680.692982	107654.047662	236.607302	4.345969	2.473684
std	9742.810427	176985.132549	515.015990	10.664949	1.736930
min	1939.000000	0.000110	0.008400	-11.920000	0.000000
25%	3336.000000	0.000878	0.107500	-6.241250	1.000000
50%	5776.000000	0.030500	0.672500	10.260000	2.000000
75%	16417.500000	198650.000000	45.250000	13.830000	4.000000
max	40000.000000	849420.000000	1948.500000	20.060000	5.000000

```

In [15]: # display the unique values of cuolumn star color
df_stars['Star color'].sort_values().unique()

```

```

Out[15]: array(['blue', 'blue white', 'red', 'white', 'white-yellow'], dtype=object)

```

```

In [16]: # format unified
df_stars['Star color'] = df_stars['Star color'].str.lower()

```

```
In [17]: # update of the similar colors name to one format
def str_replace(color):
    if color == 'blue ':
        return 'blue'
    elif color == 'blue white ' or color == 'blue-white':
        return 'blue white'
    elif color == 'yellow-white':
        return 'white-yellow'
    else:
        return color
```

```
In [18]: # deletion of implicit colors
correct_color = []
for i in df_stars['Star color']:
    correct_color.append(str_replace(i))
```

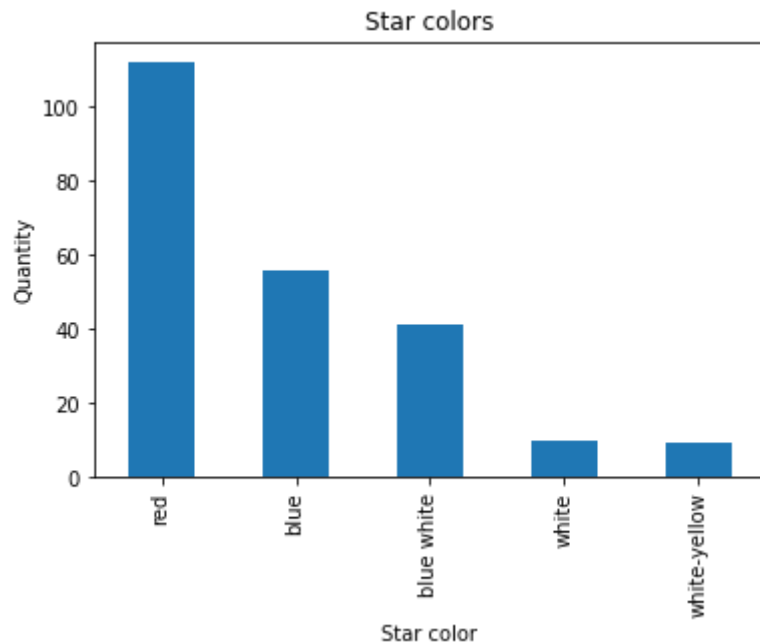
```
In [19]: df_stars['Star color'] = correct_color
```

```
In [20]: # plotting the color destrtribution

print(df_stars.groupby('Star color')['Star type'].count().sort_values(ascending=False))
df_stars.groupby('Star color')['Star type'].count().sort_values(ascending=False).plot(kind = 'bar', title = 'Star colors')
plt.xlabel('Star color')
plt.ylabel('Quantity')
```

```
Star color
red          112
blue         56
blue white   41
white        10
white-yellow  9
Name: Star type, dtype: int64
Text(0, 0.5, 'Quantity')
```

Out[20]:



```
In [21]: # selection of stars with quantity of color less than 5
star_colors = df_stars.groupby('Star color')['Star type'].count().sort_values(ascending=False).reset_index()
star_colors = star_colors.rename(columns={"Star type": "quantity"})
star_color_good = star_colors.query('quantity > 5')['Star color']
```

```
In [22]: # Deletion of colors with quantity of stars less than 5

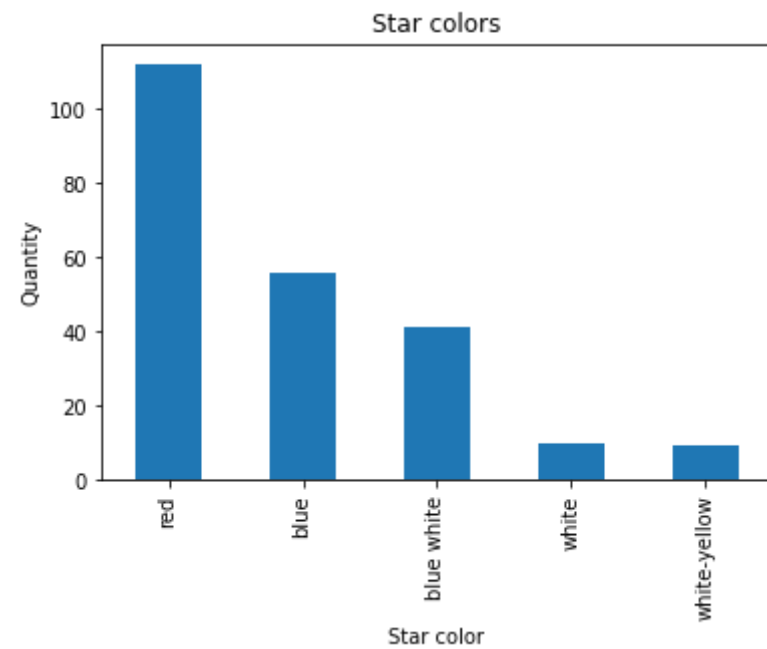
df_stars = df_stars.rename(columns={"Star color": "Star_color"})
df_stars = df_stars.query('Star_color in @star_color_good')
df_stars = df_stars.rename(columns={"Star_color": "Star color"})
```

```
In [23]: # display the color distribution table and plot it

print(df_stars.groupby('Star color')['Star type'].count().sort_values(ascending=False))
df_stars.groupby('Star color')['Star type'].count().sort_values(ascending=False).plot(kind = 'bar', title = 'Star colors')
plt.xlabel('Star color')
plt.ylabel('Quantity')
```

```
Star color
red          112
blue         56
blue white   41
white        10
white-yellow  9
Name: Star type, dtype: int64
Text(0, 0.5, 'Quantity')
```

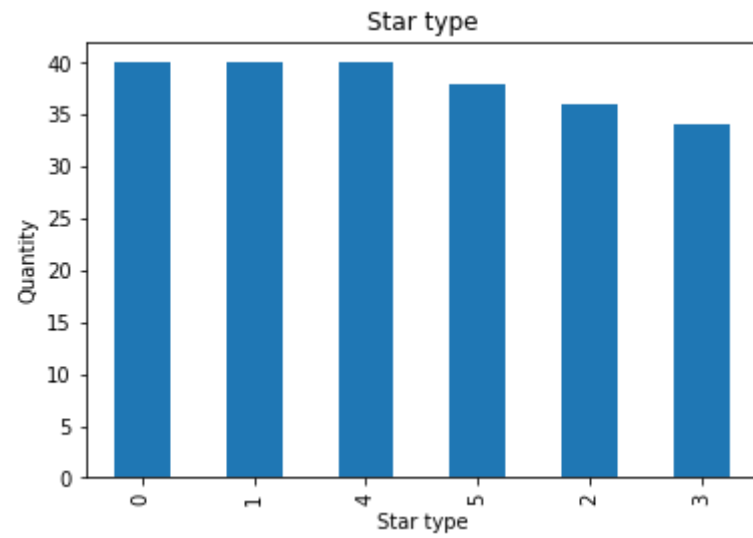
Out[23]:



In [15]: *# plotting the star type distribution*

```
df_stars.groupby('Star type')['Unnamed: 0'].count().sort_values(ascending=False).plot(kind = 'bar', title = 'Star type')
plt.xlabel('Star type')
plt.ylabel('Quantity')
```

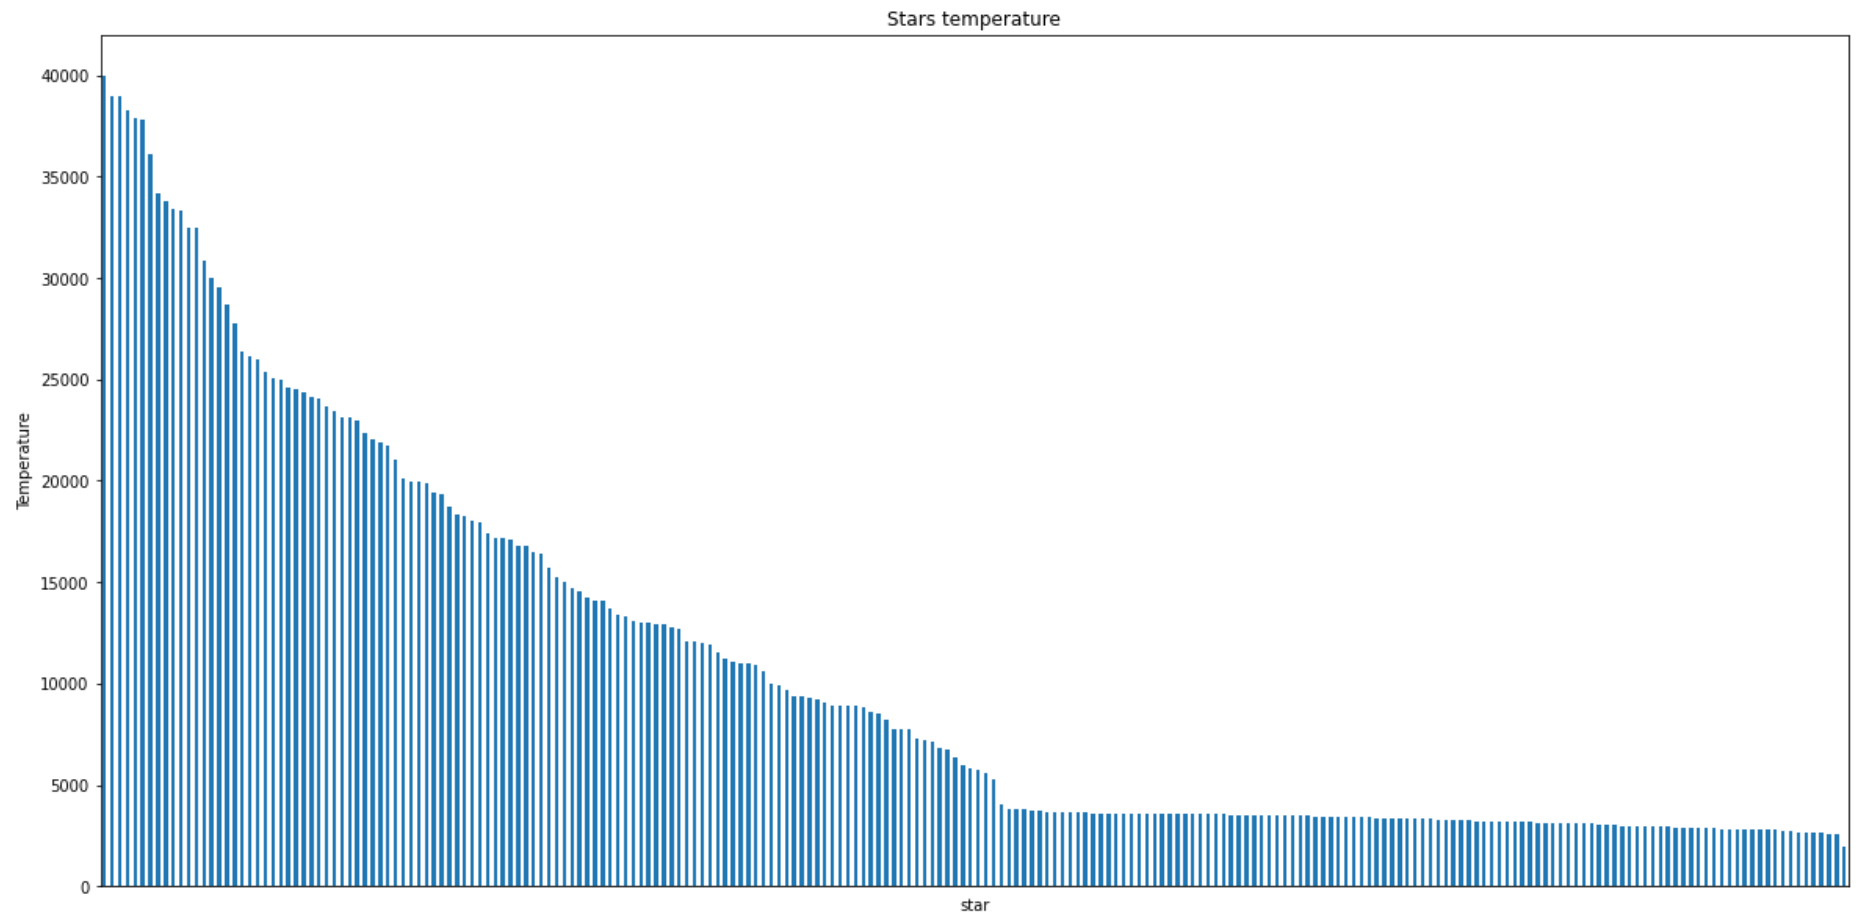
Out[15]: Text(0, 0.5, 'Quantity')



In [24]: *# plotting the stars temperature distribution*

```
df_stars['Temperature (K)'].sort_values(ascending=False).plot(kind='bar',figsize = (20,10))
plt.xlabel('star')
plt.ylabel('Temperature')
plt.title('Stars temperature')
plt.xticks([])
```

Out[24]: ([], [])

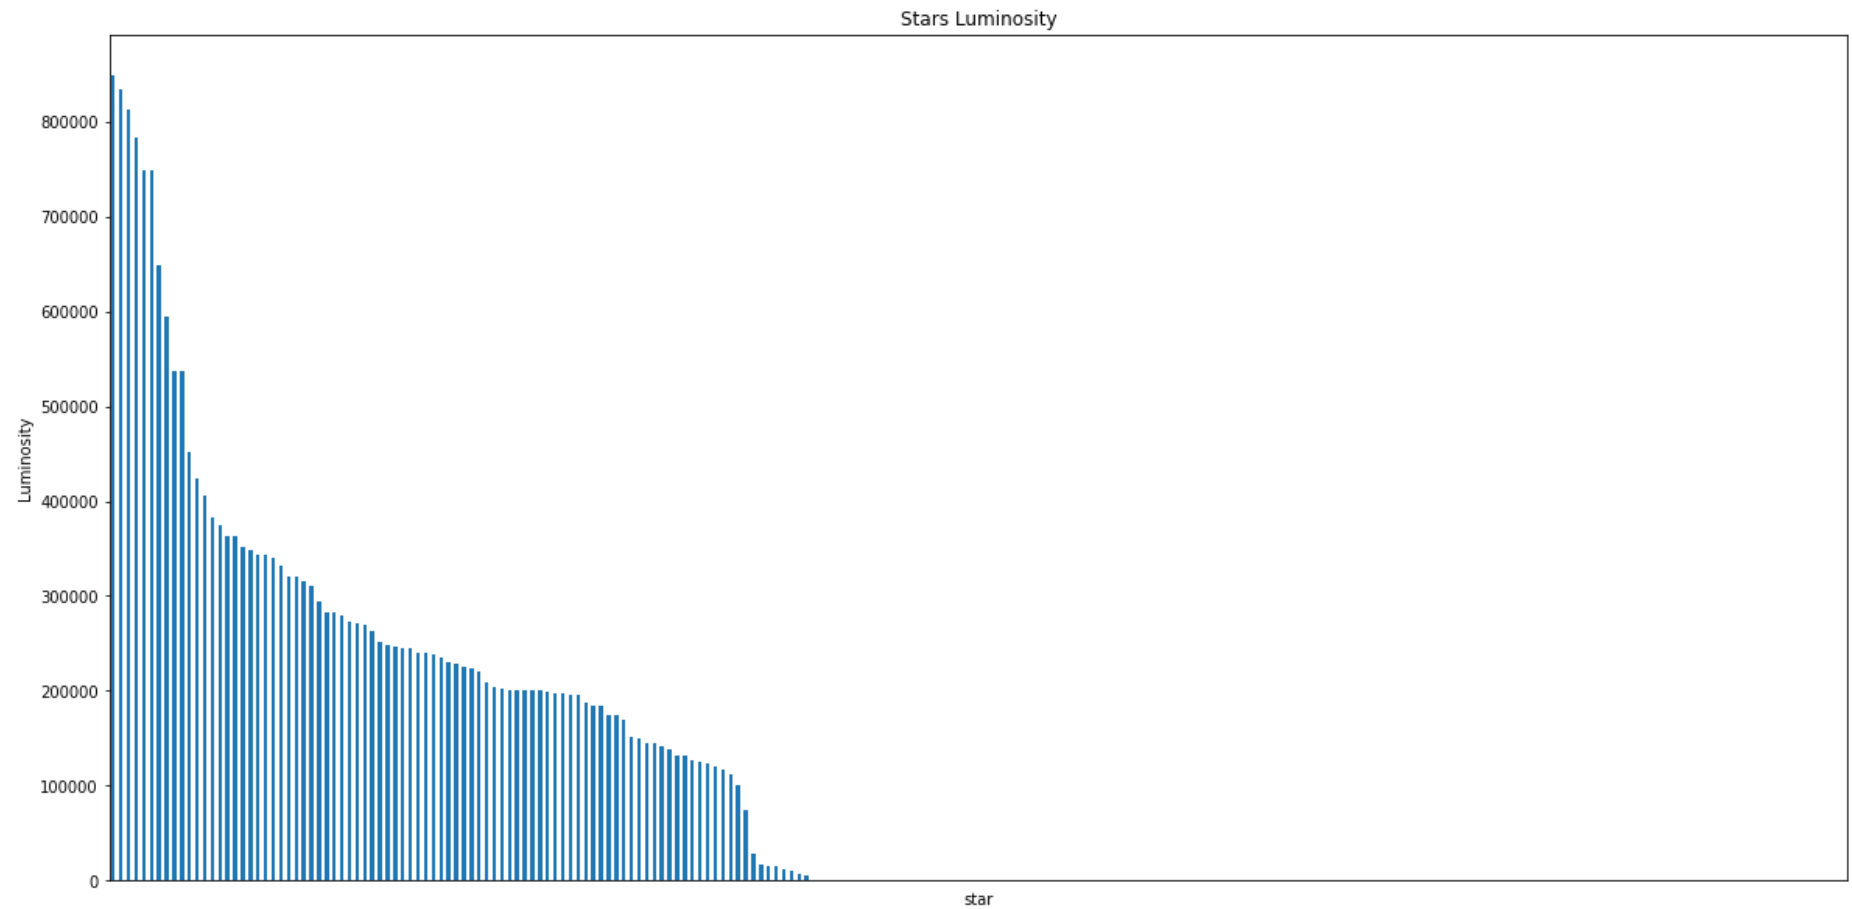


In [25]: *# plotting the stars Luminosity distribution*

```
df_stars['Luminosity(L/Lo)'].sort_values(ascending=False).plot(kind='bar',figsize = (20,10))
plt.xlabel('star')
plt.ylabel('Luminosity')
plt.title('Stars Luminosity')

plt.xticks([])
```


Out[25]: ([], [])



In [26]: *# plotting the stars radius destribution*

```
df_stars['Radius(R/Ro)'].sort_values(ascending=False).plot(kind='bar',figsize = (20,10))
plt.xlabel('star')
plt.ylabel('Radius')
plt.title('Stars Raduis')
plt.xticks([])
```

Out[26]: ([], [])



Conclusions

1) During the analysis the incorrect color names were revealed and chaned to correct.

2) Based on the information shown on histogram, the can cocnlude following:

- most part of stars color is red, blue and white-blue;
- more than a half of stars has temperature of 5000 K;
- 70 % of stars have radius below 250;
- 55% o stars has luminosity less then 40 000 units.

Data encoding

In [20]: *# separation of categorical and numeric features*

```
numeric_cols = df_stars.drop(columns = ['Star color', 'Star type']).columns  
categorical_cols = ['Star color', 'Star type']
```

```
print('Numeric:', numeric_cols, '\n', '\n', 'Categorical:', categorical_cols)
```

```
Numeric: Index(['Temperature (K)', 'Luminosity(L/Lo)', 'Radius(R/Ro)',  
              'Absolute magnitude(Mv)'],  
              dtype='object')
```

```
Categorical: ['Star color', 'Star type']
```

In [21]: *# create a copy of df*

```
df_stars_for_coding = df_stars.copy()
```

In [22]: `scaler = MinMaxScaler()`

In [23]: *# numeric features encoding*

```
scaler.fit(df_stars_for_coding[numeric_cols])  
numeric_transformed = pd.DataFrame(scaler.transform(df_stars_for_coding[numeric_cols]), columns = [numeric_cols])
```

In [24]: *# display the result*

```
numeric_transformed
```

Out[24]:

	Temperature (K)	Luminosity(L/L _o)	Radius(R/R _o)	Absolute magnitude(M _v)
--	-----------------	-------------------------------	---------------------------	-------------------------------------

0	0.029663	2.695957e-09	0.000083	0.876798
1	0.028980	4.591368e-10	0.000075	0.891807
2	0.017367	2.236820e-10	0.000048	0.957473
3	0.022622	1.059547e-10	0.000078	0.893371
4	0.000000	3.296367e-11	0.000049	1.000000
...
223	0.972150	4.412776e-01	0.695919	0.062226
224	0.759307	9.818959e-01	0.612777	0.040338
225	0.181025	6.327765e-01	0.730304	0.037211
226	0.191692	4.767253e-01	0.570694	0.021576
227	0.944352	3.471816e-01	0.915063	0.128831

228 rows × 4 columns

```
In [25]: numeric_transformed.columns = df_stars.drop(columns = ['Star color', 'Star type']).columns
```

```
In [26]: # categorical features encoding
ohe = OneHotEncoder()
categorical_transformed = ohe.fit_transform(df_stars_for_coding[categorical_cols]).toarray()
```

```
In [27]: # encoded data to dataset
categorical_transformed = pd.DataFrame(categorical_transformed)
```

```
In [28]: # display the results
categorical_transformed
```

```
Out[28]:
```

	0	1	2	3	4	5	6	7	8	9	10
0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
...
223	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
224	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
225	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
226	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
227	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0

228 rows × 11 columns

```
In [29]: # join the encoded data
df_stars_prepared = numeric_transformed.join(categorical_transformed)
```

```
In [30]: # display the results
df_stars_prepared
```

Out[30]:

	Temperature (K)	Luminosity(L/L _o)	Radius(R/R _o)	Absolute magnitude(M _v)	0	1	2	3	4	5	6	7	8	9	10
0	0.029663	2.695957e-09	0.000083	0.876798	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
1	0.028980	4.591368e-10	0.000075	0.891807	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
2	0.017367	2.236820e-10	0.000048	0.957473	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
3	0.022622	1.059547e-10	0.000078	0.893371	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
4	0.000000	3.296367e-11	0.000049	1.000000	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
...
223	0.972150	4.412776e-01	0.695919	0.062226	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
224	0.759307	9.818959e-01	0.612777	0.040338	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
225	0.181025	6.327765e-01	0.730304	0.037211	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
226	0.191692	4.767253e-01	0.570694	0.021576	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
227	0.944352	3.471816e-01	0.915063	0.128831	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0

228 rows × 15 columns

```
In [31]: # selection of target
df_stars_Y = df_stars_prepared['Temperature (K)']
```

```
In [32]: # Selection of features
df_stars_X = df_stars_prepared.drop(columns = ['Temperature (K)'])
```

```
In [33]: # split the data to train and test samples
X_train, X_test, y_train, y_test = train_test_split(df_stars_X, df_stars_Y, test_size=0.3, random_state=42)
```

```
In [34]: # data transform to tensor
X_train = torch.FloatTensor(X_train.values)
X_test = torch.FloatTensor(X_test.values)
y_train = torch.FloatTensor(y_train.values)
y_test = torch.FloatTensor(y_test.values)
```

Building Neural networks

Building base Neural network

In [35]: *# function for calculation of RMSE plotting the hist of comparison actual and predicted temperature of stars*

```
def plot_result (preds, y_test_):
    preds = preds.detach().numpy()
    y_test_ = y_test_.detach().numpy()
    df = pd.DataFrame(preds, columns = ['prediction'])
    df['actual'] = y_test_
    for i in df.columns:
        df[i] = df[i] * df_stars['Temperature (K)'].max()
    df.plot(kind = 'bar',figsize=(20,10))
    plt.title('star temperature')
    plt.ylabel('temprature')
    plt.xticks([])
    print(df.head())
    print('\n', 'RMSE:', mean_squared_error(df['actual'],df['prediction'],squared = False))
```

In [36]: *# Net 1 building*

```
torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8
output_size = 1

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.fc3 = nn.Linear(hidden_size_2, output_size)

    def forward(self, x):
        x = self.fc1(x)
        x = self.fc2(x)
        x = self.fc3(x)
        return x
```

```
net_1 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)
```

In [37]: *# net 1 training and making a prediction*

```
optimizer = torch.optim.Adam(net_1.parameters(),lr=0.001)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_1.forward(X_train).flatten()
    loss_value = loss(preds,y_train)
    loss_value.backward()
    optimizer.step()
    if (epoch % 25 == 0) or (epoch == 500) :
        print(loss_value)
        net_1.eval(),
        test_preds_1 = net_1.forward(X_test).flatten()
        accuracy = (torch.round(test_preds_1) == y_test).float().mean().data
        print(accuracy)
```




```
tensor(0.1294, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0775, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0485, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0294, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0226, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0195, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0176, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0166, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0163, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0162, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0161, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0161, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0161, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0161, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0161, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)
```

```
In [38]: # net results  
plot_result(test_preds_1,y_test)
```

	prediction	actual
0	25638.511719	9449.041992
1	1454.730591	1272.693848
2	1245.530518	799.768799
3	17902.568359	11718.031250
4	1319.206055	1472.373291

RMSE: 4868.386



Вывод по сети_1

RMSE для сети_1 приемлемое можно использовать данную сеть для ее дальнейшего улучшения. Хотя сеть иногда дает отрицательные значения...

Building neural network with LeakyReLU activation

In [39]: *# bulding net_2*

```
torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8
output_size = 1

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.act1 = nn.LeakyReLU()
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.act2 = nn.LeakyReLU()
        self.fc3 = nn.Linear(hidden_size_2, output_size)
        self.act3 = nn.LeakyReLU()

    def forward(self, x):
        x = self.fc1(x)
        x = self.act1(x)
        x = self.fc2(x)
        x = self.act2(x)
        x = self.fc3(x)
        x = self.act3(x)
        return x

net_2 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)
```

In [40]: *# net_2 training and making a prediction*

```
optimizer = torch.optim.Adam(net_2.parameters(), lr=0.001)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
    optimizer.zero_grad()
```


```
preds = net_2.forward(X_train).flatten()
loss_value = loss(preds,y_train)
loss_value.backward()
optimizer.step()
if (epoch % 25 == 0) or (epoch == 500) :
    print(loss_value)
    net_2.eval(),
    test_preds_2 = net_2.forward(X_test).flatten()
    accuracy = (torch.round(test_preds_2) == y_test).float().mean().data
    print(accuracy)
```

```
tensor(0.1117, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0811, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0615, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0446, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0314, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0250, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0218, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0194, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0172, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0154, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0142, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0135, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0131, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0128, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0126, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0124, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0122, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0120, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0118, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0116, grad_fn=<MseLossBackward0>)  
tensor(0.)
```

```
In [41]: plot_result(test_preds_2,y_test)
```

	prediction	actual
0	21669.697266	9449.041992
1	1492.615234	1272.693848
2	1537.019897	799.768799
3	14472.894531	11718.031250
4	1318.640869	1472.373291

RMSE: 4356.7583



Вывод по сети_2

RMSE для сети_2 уже достигло необходимых показателей, но зачем останавливаться если ее можно использовать для дальнейшего улучшения.

Building neural network with Sigmoid activation function

```
In [42]: # building net_3

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.act1 = nn.Sigmoid()
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.act2 = nn.Sigmoid()
        self.fc3 = nn.Linear(hidden_size_2, output_size)
        self.act3 = nn.Sigmoid()

    def forward(self, x):
        x = self.fc1(x)
        x = self.act1(x)
        x = self.fc2(x)
        x = self.act2(x)
        x = self.fc3(x)
        x = self.act3(x)
        return x

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 18
hidden_size_2 = 45
output_size = 1

net_3 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)
```

```
In [43]: # net_3 training and making a prediction

optimizer = torch.optim.Adam(net_3.parameters(), lr=0.001)

loss = torch.nn.MSELoss()

num_epochs = 600

for epoch in range(num_epochs):
```

```

optimizer.zero_grad()
preds = net_2.forward(X_train).flatten()
loss_value = loss(preds,y_train)
loss_value.backward()
optimizer.step()
if (epoch % 100 == 0) or (epoch == 600) :
    net_3.eval(),
    test_preds_3 = net_3.forward(X_test).flatten()
    accuracy = (torch.round(test_preds_3) == y_test).float().mean().data
    print(accuracy)

```

```

tensor(0.)
tensor(0.)
tensor(0.)
tensor(0.)
tensor(0.)
tensor(0.)

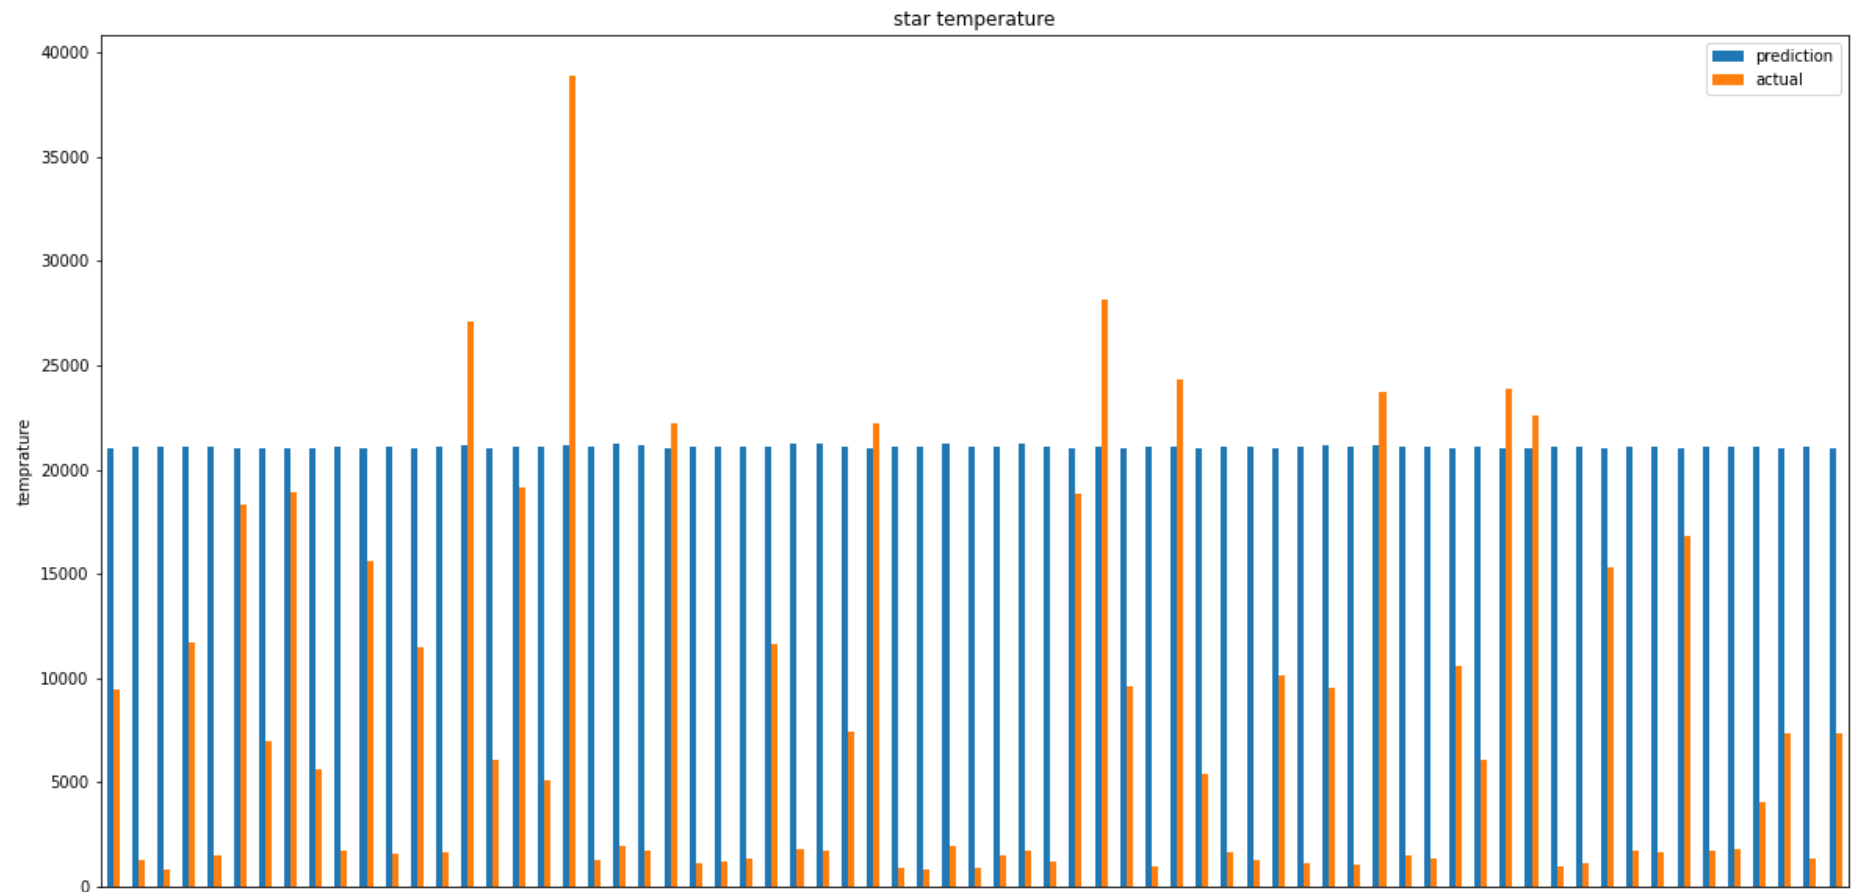
```

In [44]: *# display the results of net_3*

```
plot_result(test_preds_3,y_test)
```

	prediction	actual
0	21032.304688	9449.041992
1	21063.041016	1272.693848
2	21090.667969	799.768799
3	21094.648438	11718.031250
4	21062.929688	1472.373291

RMSE: 15818.15



Вывод по сети_3

Функция сигмоид дает среднее значение - такую нейронную сеть предлагаю не использовать для улучшения параметров

Neural Networks tuning

learning rate tuning

```
In [45]: # based on net 1 building net_4
```

```

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8
output_size = 1

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.fc3 = nn.Linear(hidden_size_2, output_size)

    def forward(self, x):
        x = self.fc1(x)
        x = self.fc2(x)
        x = self.fc3(x)
        return x

net_4 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)

```

In [46]: *# net 4 training and making a prediction (changes are - lr +=0.01)*

```

optimizer = torch.optim.Adam(net_4.parameters(),lr=0.011)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_4.forward(X_train).flatten()
    loss_value = loss(preds,y_train)
    loss_value.backward()
    optimizer.step()
    if (epoch % 25 == 0) or (epoch == 500) :
        print(loss_value)
        net_4.eval(),
        test_preds_4 = net_4.forward(X_test).flatten()
        accuracy = (torch.round(test_preds_4) == y_test).float().mean().data
        print(accuracy)

```

```

tensor(0.1294, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0166, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0162, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0160, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0159, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0159, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0158, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0158, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0157, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0157, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0156, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0156, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0156, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0156, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0156, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0156, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0156, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0156, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0156, grad_fn=<MseLossBackward0>)
tensor(0.)

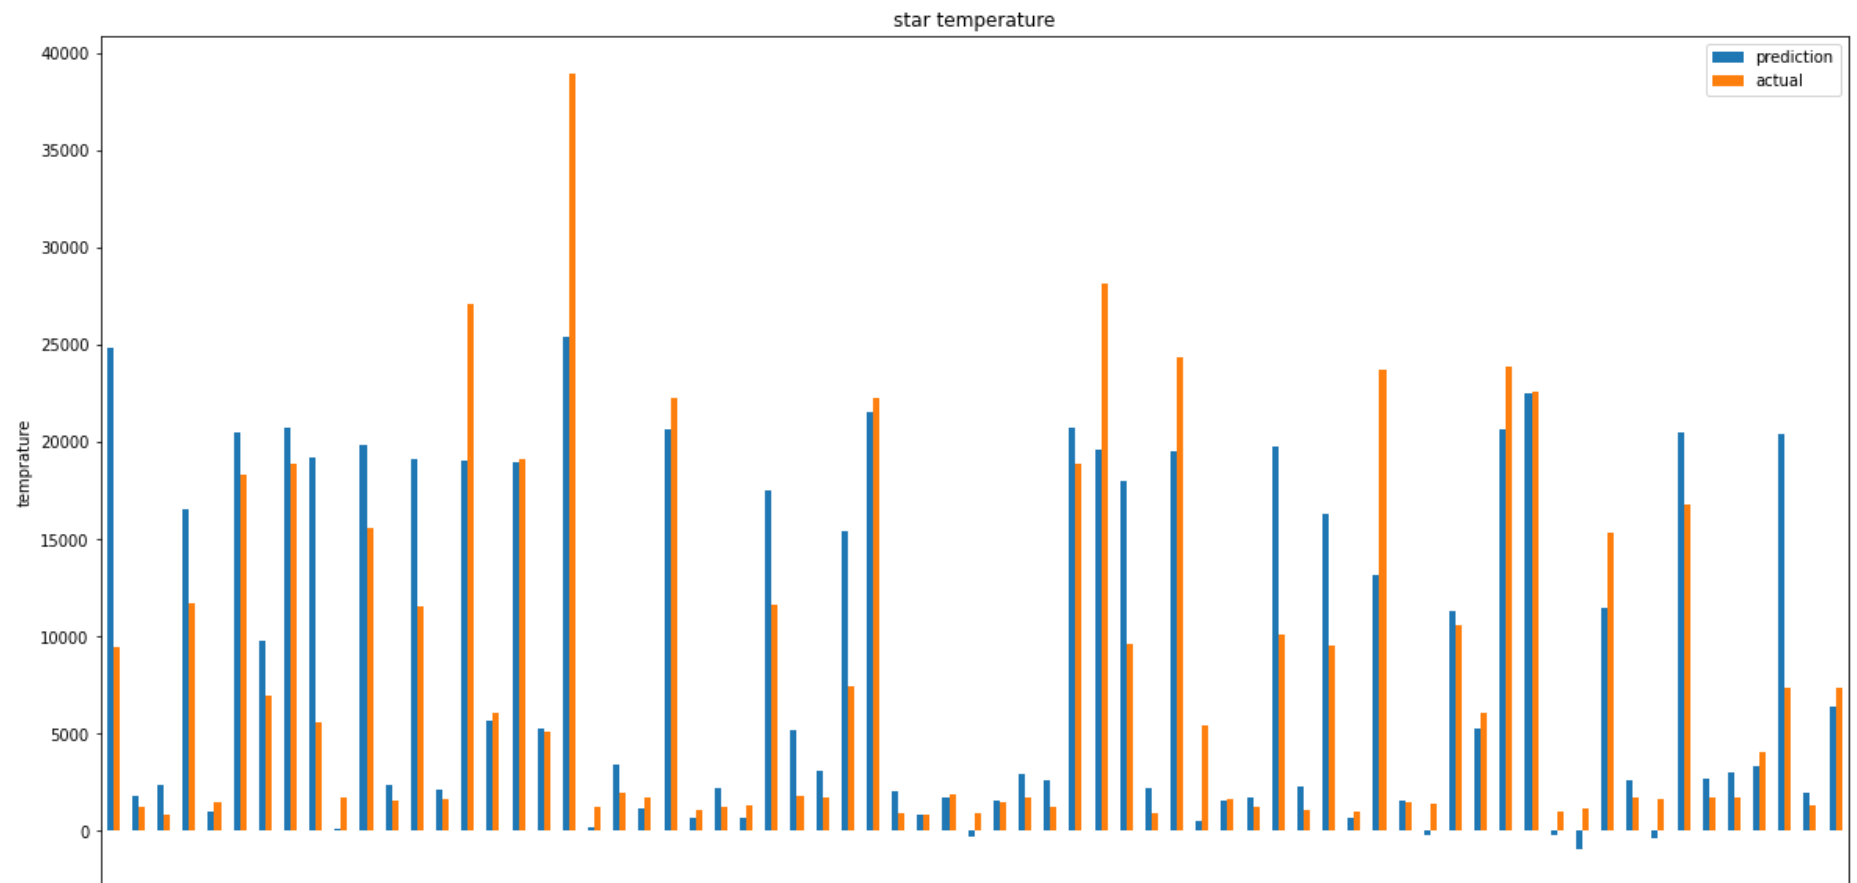
```

In [47]: *# display the results*

```
plot_result(test_preds_4,y_test)
```

	prediction	actual
0	24865.097656	9449.041992
1	1804.313721	1272.693848
2	2341.796387	799.768799
3	16544.482422	11718.031250
4	977.397095	1472.373291

RMSE: 4803.809



In [48]: *# based on net 2 building the net_5*

```

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8

```

```

output_size = 1

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.act1 = nn.LeakyReLU()
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.act2 = nn.LeakyReLU()
        self.fc3 = nn.Linear(hidden_size_2, output_size)
        self.act3 = nn.LeakyReLU()

    def forward(self, x):
        x = self.fc1(x)
        x = self.act1(x)
        x = self.fc2(x)
        x = self.act2(x)
        x = self.fc3(x)
        x = self.act3(x)
        return x

net_5 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)

```

```

In [49]: # net_5 training and making a prediction (changes are - lr +=0.001)

optimizer = torch.optim.Adam(net_5.parameters(),lr=0.002)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_5.forward(X_train).flatten()
    loss_value = loss(preds,y_train)
    loss_value.backward()
    optimizer.step()
    if (epoch % 25 == 0) or (epoch == 500) :
        print(loss_value)
        net_5.eval(),
        test_preds_5 = net_5.forward(X_test).flatten()
        accuracy = (torch.round(test_preds_5) == y_test).float().mean().data
        print(accuracy)

```

```

tensor(0.1117, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0628, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0349, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0232, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0193, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0167, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0150, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0142, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0136, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0132, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0128, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0123, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0118, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0111, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0107, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0103, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0100, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0098, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0097, grad_fn=<MseLossBackward0>)
tensor(0.)
tensor(0.0096, grad_fn=<MseLossBackward0>)
tensor(0.)


```

In [50]: *# displaay the results*

```
plot_result(test_preds_5,y_test)
```

	prediction	actual
0	21693.945312	9449.041992
1	-8.832759	1272.693848
2	-15.238729	799.768799
3	9485.941406	11718.031250
4	-8.603992	1472.373291

RMSE: 4400.8994



Learning rate changes

- Changes of learning rate has slightly increase the RMSE score of net 1 - changes are accepted and to be used for the further model optimisation.
- Changes of learning rate has reduce the RMSE score of net 2 - changes rejected.

Optimizer tuning

```
In [52]: # based on optimised net 1 building the net_1

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8
output_size = 1

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.fc3 = nn.Linear(hidden_size_2, output_size)

    def forward(self, x):
        x = self.fc1(x)
        x = self.fc2(x)
        x = self.fc3(x)
        return x

net_4 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)
```

```
In [53]: # training of net_4 (changes - optimizer - SGD)

optimizer = torch.optim.SGD(net_4.parameters(), lr=0.011)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_4.forward(X_train).flatten()
    loss_value = loss(preds, y_train)
    loss_value.backward()
    optimizer.step()
    if (epoch % 25 == 0) or (epoch == 500) :
        print(loss_value)
```



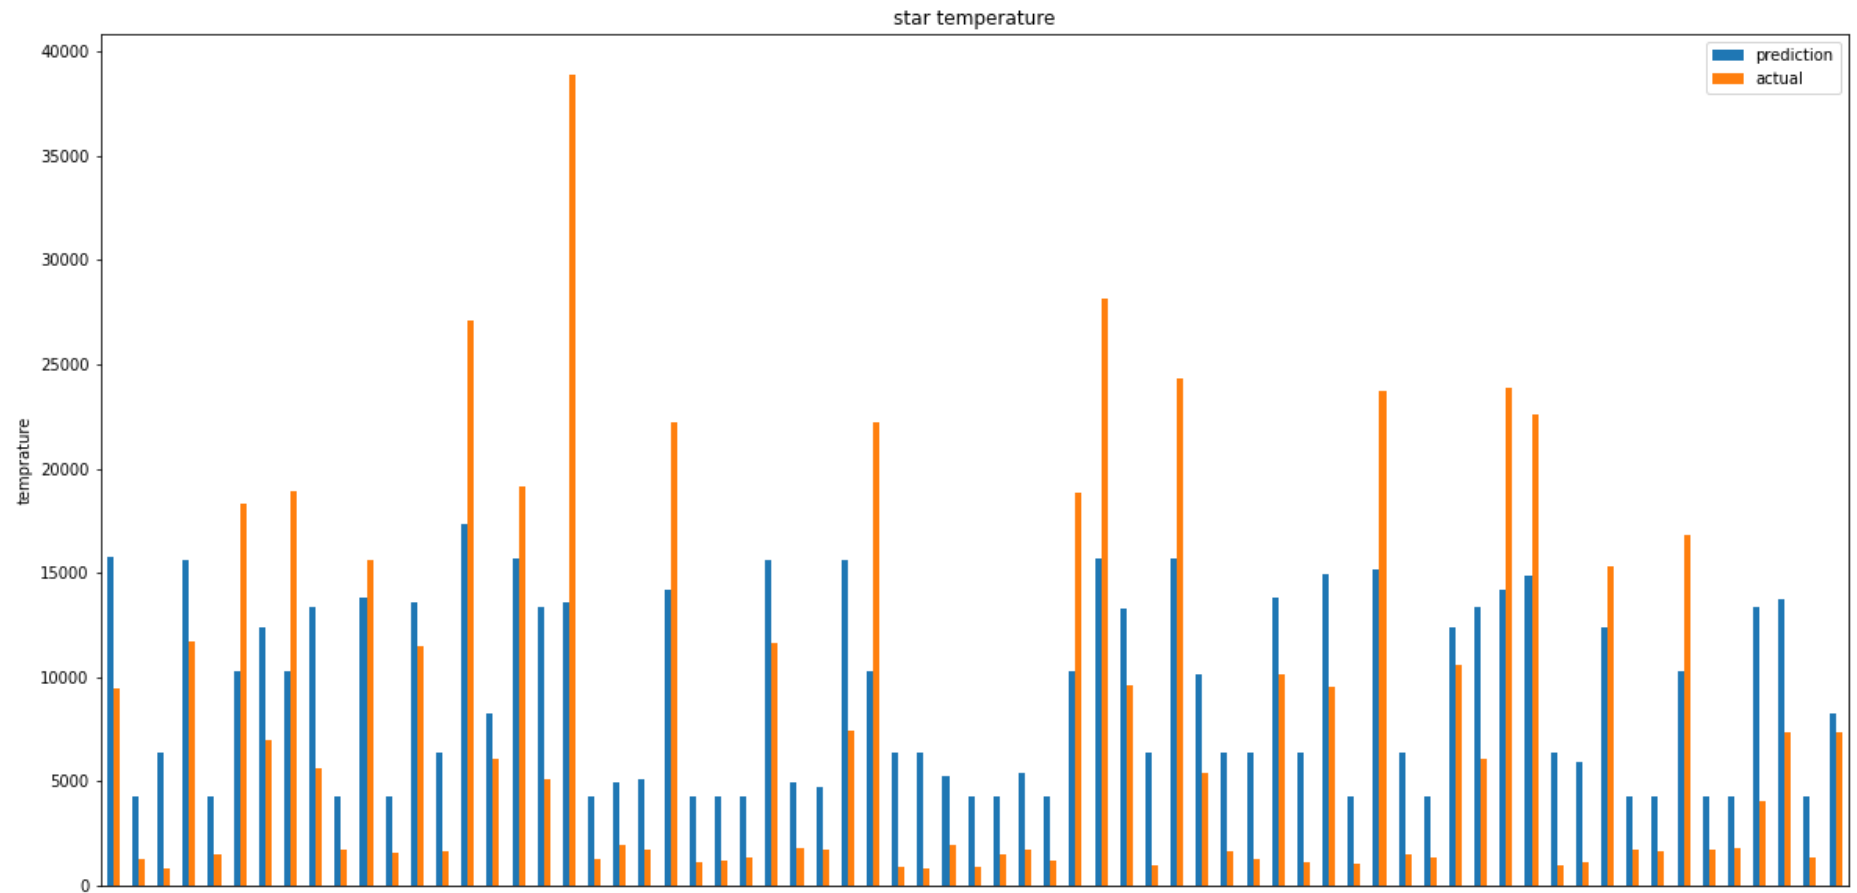
```
net_4.eval(),
test_preds_4 = net_4.forward(X_test).flatten()
accuracy = (torch.round(test_preds_4) == y_test).float().mean().data
print(accuracy)
```

```
tensor(0.1294, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0986, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0893, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0840, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0798, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0763, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0731, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0702, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0674, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0648, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0622, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0596, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0570, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0545, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0519, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0494, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0469, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0444, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0420, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0397, grad_fn=<MseLossBackward0>)  
tensor(0.)
```

```
In [54]: plot_result(test_preds_4,y_test)
```

	prediction	actual
0	15725.837891	9449.041992
1	4302.927246	1272.693848
2	6401.625977	799.768799
3	15632.453125	11718.031250
4	4302.507324	1472.373291

RMSE: 6321.381



In [55]: *# based on optimised net 1 building the net_4*

```

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8

```

```

output_size = 1

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.fc3 = nn.Linear(hidden_size_2, output_size)

    def forward(self, x):
        x = self.fc1(x)
        x = self.fc2(x)
        x = self.fc3(x)
        return x

net_4 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)

```

```

In [56]: # training of net_4 (changes - optimizer - RMSprop)

optimizer = torch.optim.RMSprop(net_4.parameters(), lr=0.001)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_4.forward(X_train).flatten()
    loss_value = loss(preds, y_train)
    loss_value.backward()
    optimizer.step()
    if (epoch % 25 == 0) or (epoch == 500) :
        print(loss_value)
        net_4.eval(),
        test_preds_4 = net_4.forward(X_test).flatten()
        accuracy = (torch.round(test_preds_4) == y_test).float().mean().data
        print(accuracy)

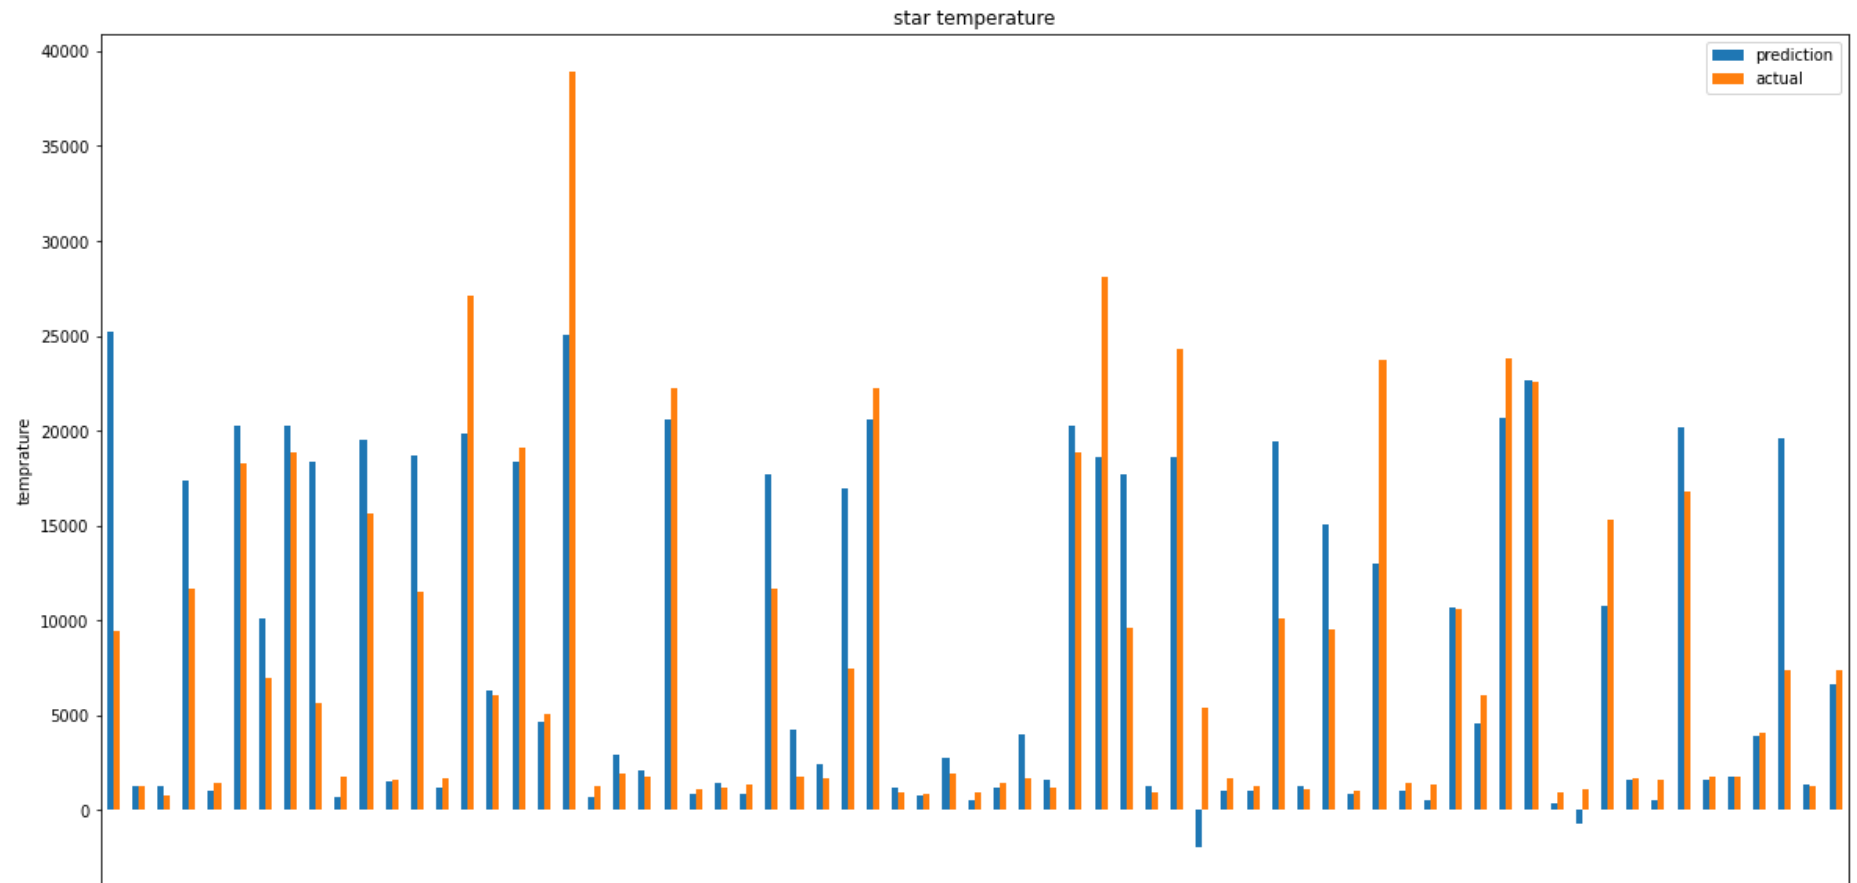
```

```
tensor(0.1294, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0239, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0180, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0165, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0162, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0161, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0161, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0158, grad_fn=<MseLossBackward0>)  
tensor(0.)
```

```
In [57]: plot_result(test_preds_4,y_test)
```

	prediction	actual
0	25195.689453	9449.041992
1	1302.915405	1272.693848
2	1295.320801	799.768799
3	17401.923828	11718.031250
4	998.719788	1472.373291

RMSE: 4810.8184



In [58]: *# based on net 2 building the net_5*

```

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8
output_size = 1

```

```

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.act1 = nn.LeakyReLU()
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.act2 = nn.LeakyReLU()
        self.fc3 = nn.Linear(hidden_size_2, output_size)
        self.act3 = nn.LeakyReLU()

    def forward(self, x):
        x = self.fc1(x)
        x = self.act1(x)
        x = self.fc2(x)
        x = self.act2(x)
        x = self.fc3(x)
        x = self.act3(x)
        return x

net_5 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)

```

```

In [59]: # training net_5 (changes - optimizer - SGD)

optimizer = torch.optim.SGD(net_5.parameters(), lr=0.001)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_5.forward(X_train).flatten()
    loss_value = loss(preds, y_train)
    loss_value.backward()
    optimizer.step()
    if (epoch % 25 == 0) or (epoch == 500) :
        print(loss_value)
        net_5.eval(),
        test_preds_5 = net_5.forward(X_test).flatten()
        accuracy = (torch.round(test_preds_5) == y_test).float().mean().data
        print(accuracy)

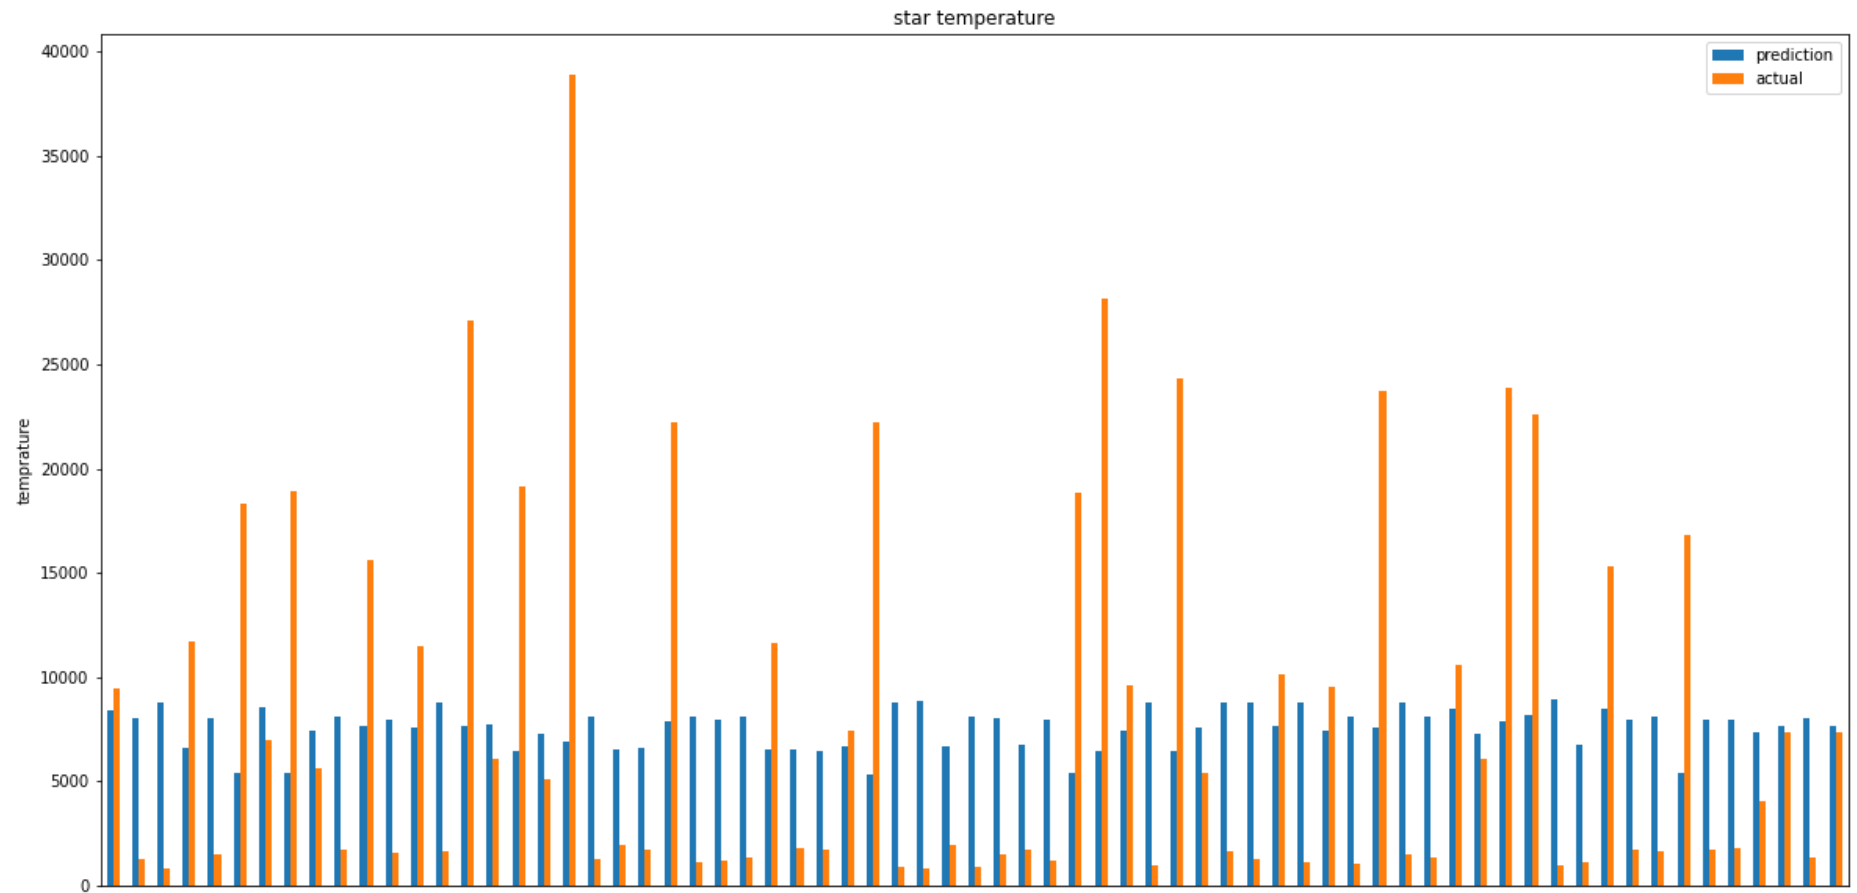
```

```
tensor(0.1117, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.1076, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.1040, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.1008, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0980, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0955, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0933, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0914, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0897, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0882, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0869, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0857, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0846, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0837, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0829, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0822, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0815, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0809, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0804, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0799, grad_fn=<MseLossBackward0>)  
tensor(0.)
```

```
In [60]: plot_result(test_preds_5,y_test)
```


	prediction	actual
0	8418.826172	9449.041992
1	7992.934082	1272.693848
2	8766.971680	799.768799
3	6583.833496	11718.031250
4	8045.629395	1472.373291

RMSE: 9329.265



In [61]: *# based on net 2 building the net_5*

```

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8

```

```

output_size = 1

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.act1 = nn.LeakyReLU()
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.act2 = nn.LeakyReLU()
        self.fc3 = nn.Linear(hidden_size_2, output_size)
        self.act3 = nn.LeakyReLU()

    def forward(self, x):
        x = self.fc1(x)
        x = self.act1(x)
        x = self.fc2(x)
        x = self.act2(x)
        x = self.fc3(x)
        x = self.act3(x)
        return x

net_5 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)

```

```

In [62]: # training net5 (changes - optimizer - RMSprop)

optimizer = torch.optim.RMSprop(net_5.parameters(),lr=0.001)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_5.forward(X_train).flatten()
    loss_value = loss(preds,y_train)
    loss_value.backward()
    optimizer.step()
    if (epoch % 25 == 0) or (epoch == 500) :
        print(loss_value)
        net_5.eval(),
        test_preds_5 = net_5.forward(X_test).flatten()
        accuracy = (torch.round(test_preds_5) == y_test).float().mean().data
        print(accuracy)

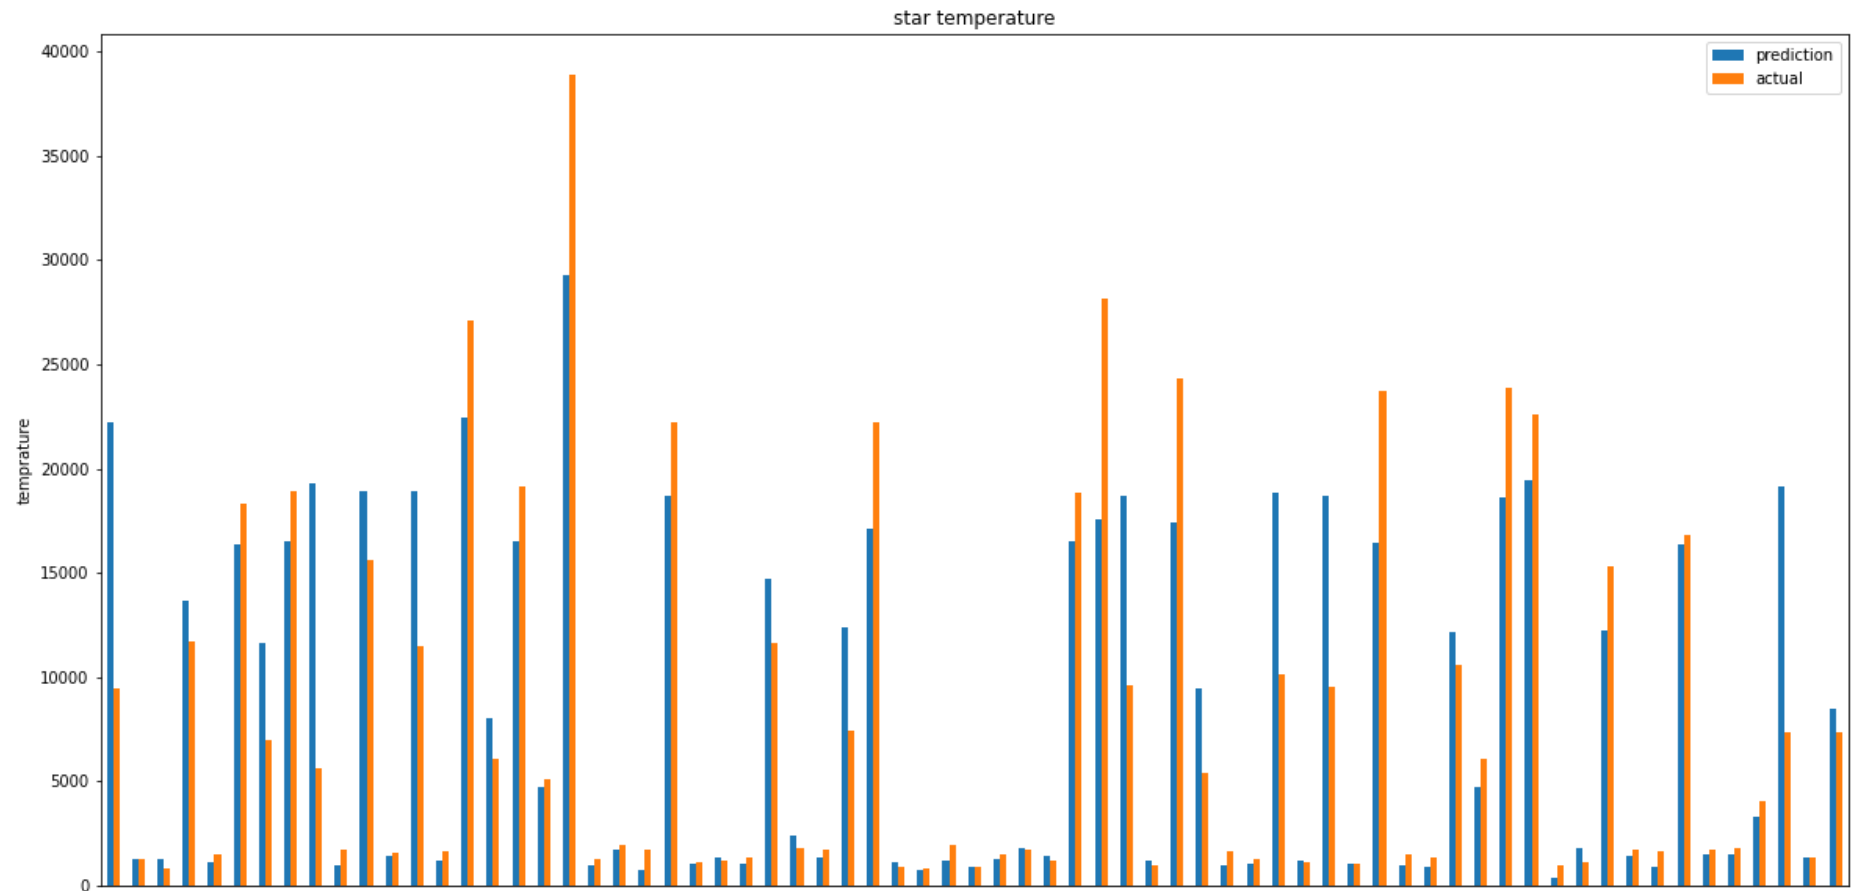
```

```
tensor(0.1117, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0342, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0228, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0181, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0156, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0141, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0132, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0129, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0127, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0126, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0125, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0124, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0122, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0121, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0120, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0118, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0117, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0115, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0114, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0112, grad_fn=<MseLossBackward0>)  
tensor(0.)
```

```
In [63]: plot_result(test_preds_5,y_test)
```

	prediction	actual
0	22199.255859	9449.041992
1	1284.554321	1272.693848
2	1227.034912	799.768799
3	13621.442383	11718.031250
4	1125.555298	1472.373291

RMSE: 4377.1074



Conclusion on optimiser changes

Changes of optimiser has not increase the RMSE score on both networks - adam optimiser will remain as the best for the networks and used for further optimisation.

Dropout tuning

```
In [64]: # based on optimised net 1 building the net_4

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8
output_size = 1

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.dp1 = nn.Dropout(p = 0.2)
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        # self.dp2 = nn.Dropout(p = 0.2)
        self.fc3 = nn.Linear(hidden_size_2, output_size)
        # self.dp3 = nn.Dropout(p = 0.2)

    def forward(self, x):
        x = self.fc1(x)
        x = self.dp1(x)
        x = self.fc2(x)
        # x = self.dp2(x)
        x = self.fc3(x)
        # x = self.dp3(x)
        return x

net_4 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)
```

```
In [65]: # training net_4 (changes - added dp1, p = 0.2)

optimizer = torch.optim.Adam(net_4.parameters(), lr=0.011)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_4.forward(X_train).flatten()
```


```
loss_value = loss(preds,y_train)
loss_value.backward()
optimizer.step()
if (epoch % 25 == 0) or (epoch == 500) :
    print(loss_value)
    net_4.eval(),
    test_preds_4 = net_4.forward(X_test).flatten()
    accuracy = (torch.round(test_preds_4) == y_test).float().mean().data
    print(accuracy)
```

[illegible]

```
In [66]: plot_result(test_preds_4,y_test)
```

	prediction	actual
0	24865.228516	9449.041992
1	1804.252808	1272.693848
2	2341.591797	799.768799
3	16544.748047	11718.031250
4	977.463867	1472.373291

RMSE: 4803.806



In [67]: *# based on optimised net 1 building the net_4*

```

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8
output_size = 1

```



```

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.dp1 = nn.Dropout(p = 0.2)
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.dp2 = nn.Dropout(p = 0.5)
        self.fc3 = nn.Linear(hidden_size_2, output_size)
        # self.dp3 = nn.Dropout(p = 0.2)

    def forward(self, x):
        x = self.fc1(x)
        x = self.dp1(x)
        x = self.fc2(x)
        x = self.dp2(x)
        x = self.fc3(x)
        # x = self.dp3(x)
        return x

net_4 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)

```

```

In [68]: # training net_4 (changes - added dp1, p = 0.2 u dp2, p = 0.5)

optimizer = torch.optim.Adam(net_4.parameters(), lr=0.011)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_4.forward(X_train).flatten()
    loss_value = loss(preds, y_train)
    loss_value.backward()
    optimizer.step()
    if (epoch % 25 == 0) or (epoch == 500) :
        print(loss_value)
        net_4.eval(),
        test_preds_4 = net_4.forward(X_test).flatten()
        accuracy = (torch.round(test_preds_4) == y_test).float().mean().data
        print(accuracy)

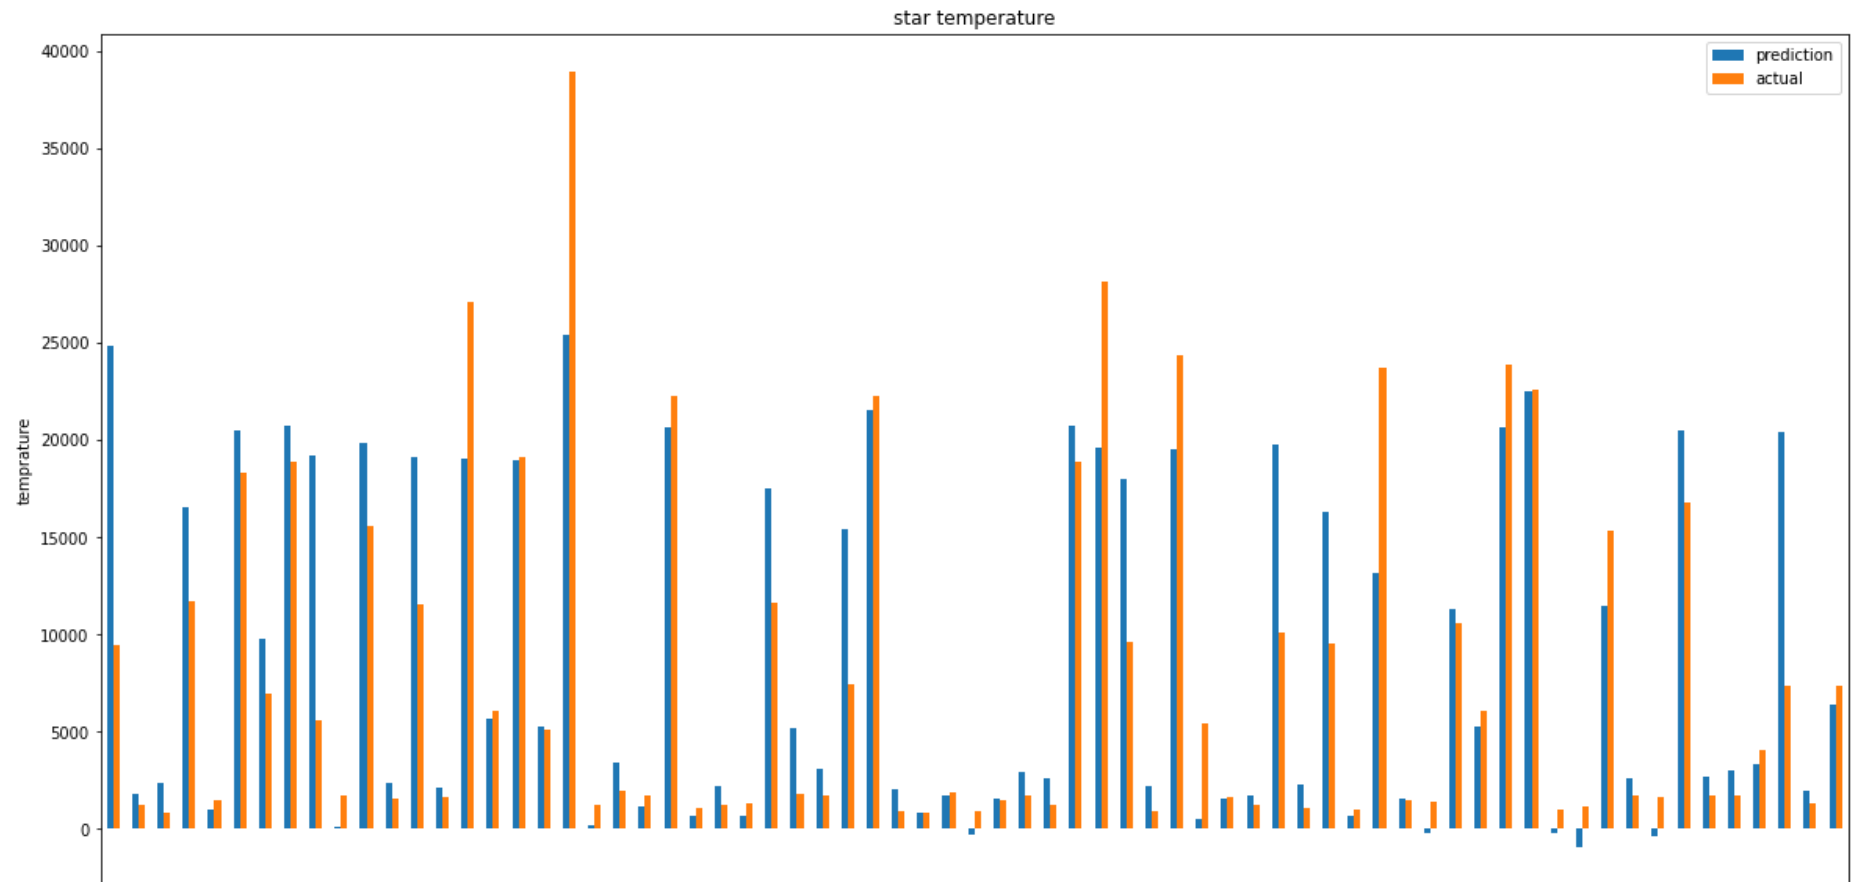
```

[illegible]

```
In [69]: plot_result(test_preds_4,y_test)
```

	prediction	actual
0	24865.097656	9449.041992
1	1804.310059	1272.693848
2	2341.790039	799.768799
3	16544.482422	11718.031250
4	977.392517	1472.373291

RMSE: 4803.809



In [70]: *# based on optimised net 1 building the net_4*

```

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8
output_size = 1

```

```

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.dp1 = nn.Dropout(p = 0.2)
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        # self.dp2 = nn.Dropout(p = 0.2)
        self.fc3 = nn.Linear(hidden_size_2, output_size)
        self.dp3 = nn.Dropout(p = 0.5)

    def forward(self, x):
        x = self.fc1(x)
        x = self.dp1(x)
        x = self.fc2(x)
        # x = self.dp2(x)
        x = self.fc3(x)
        x = self.dp3(x)
        return x

net_4 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)

```

```

In [71]: # training net_4 (changes - added dp1 u dp3, p = 0.5)

optimizer = torch.optim.Adam(net_4.parameters(),lr=0.011)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_4.forward(X_train).flatten()
    loss_value = loss(preds,y_train)
    loss_value.backward()
    optimizer.step()
    if (epoch % 25 == 0) or (epoch == 500) :
        print(loss_value)
        net_4.eval(),
        test_preds_4 = net_4.forward(X_test).flatten()
        accuracy = (torch.round(test_preds_4) == y_test).float().mean().data
        print(accuracy)

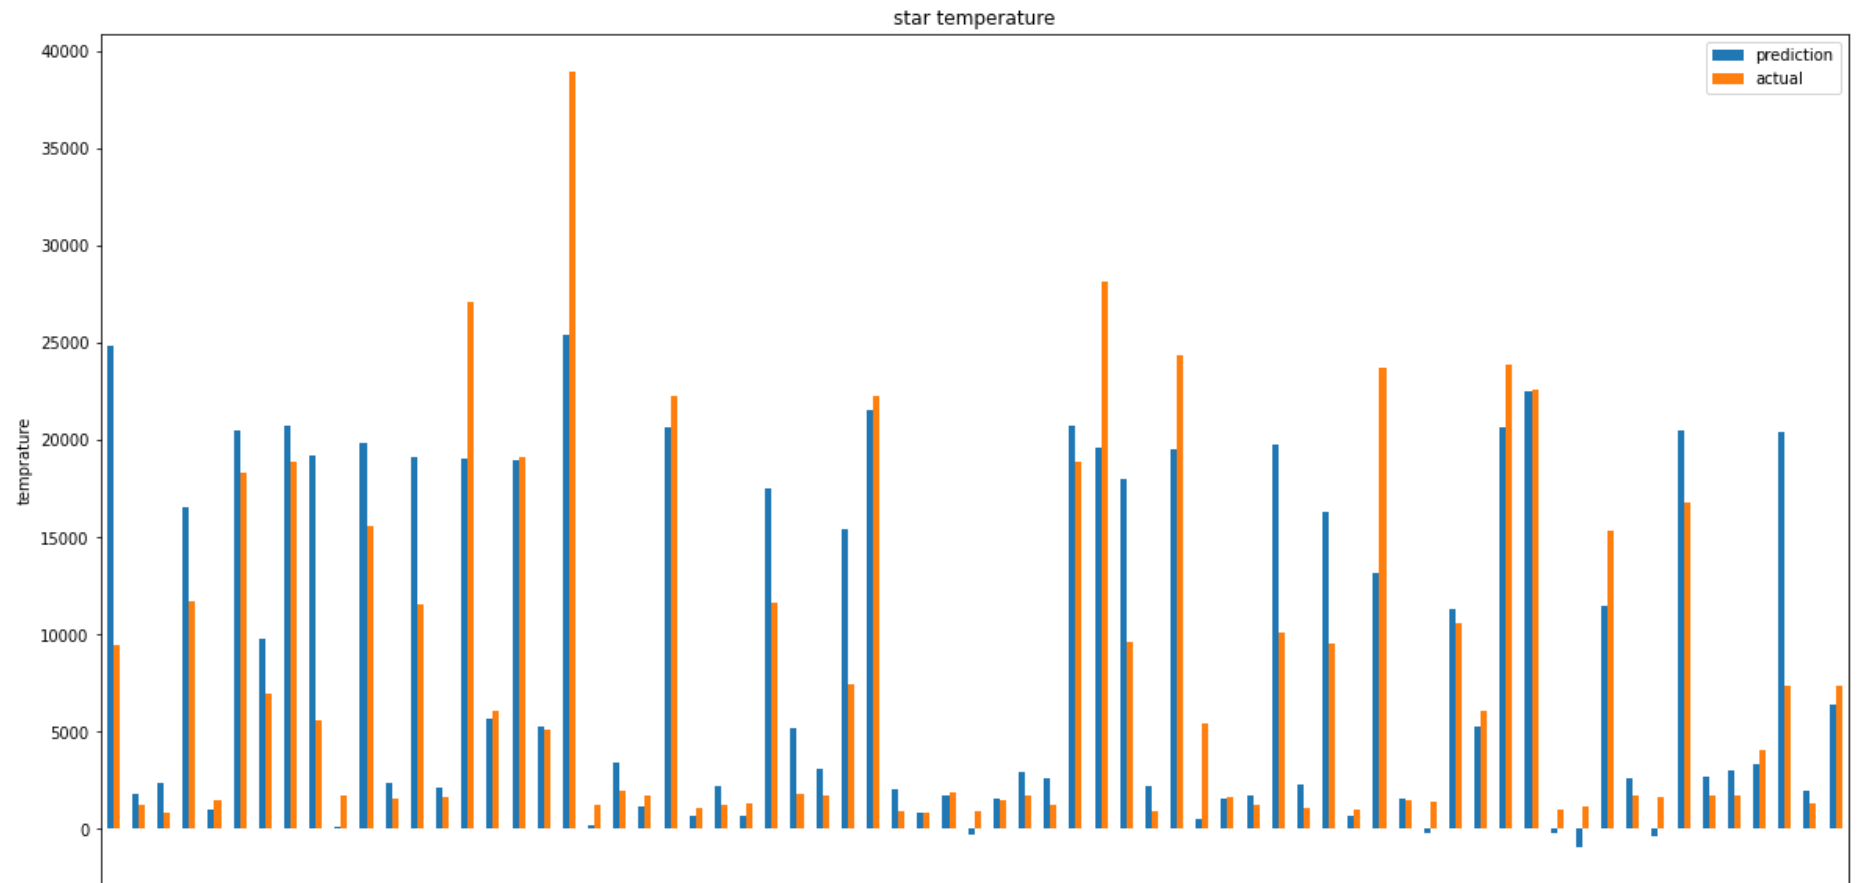
```

[illegible]

```
In [72]: plot_result(test_preds_4,y_test)
```

	prediction	actual
0	24864.966797	9449.041992
1	1804.372925	1272.693848
2	2341.990723	799.768799
3	16544.232422	11718.031250
4	977.327820	1472.373291

RMSE: 4803.8115



In [73]: *# based on optimised net 1 building the net_4*

```

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8
output_size = 1

```

```

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.dp1 = nn.Dropout(p = 0.2)
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.dp2 = nn.Dropout(p = 0.5)
        self.fc3 = nn.Linear(hidden_size_2, output_size)
        self.dp3 = nn.Dropout(p = 0.2)

    def forward(self, x):
        x = self.fc1(x)
        x = self.dp1(x)
        x = self.fc2(x)
        x = self.dp2(x)
        x = self.fc3(x)
        x = self.dp3(x)
        return x

net_4 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)

```

```

In [74]: # training net_4 (changes - added dp1, dp2 u dp3, p = 0.2, 0.5, 0.2)

optimizer = torch.optim.Adam(net_4.parameters(),lr=0.011)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_4.forward(X_train).flatten()
    loss_value = loss(preds,y_train)
    loss_value.backward()
    optimizer.step()
    if (epoch % 25 == 0) or (epoch == 500) :
        print(loss_value)
        net_4.eval(),
        test_preds_4 = net_4.forward(X_test).flatten()
        accuracy = (torch.round(test_preds_4) == y_test).float().mean().data
        print(accuracy)

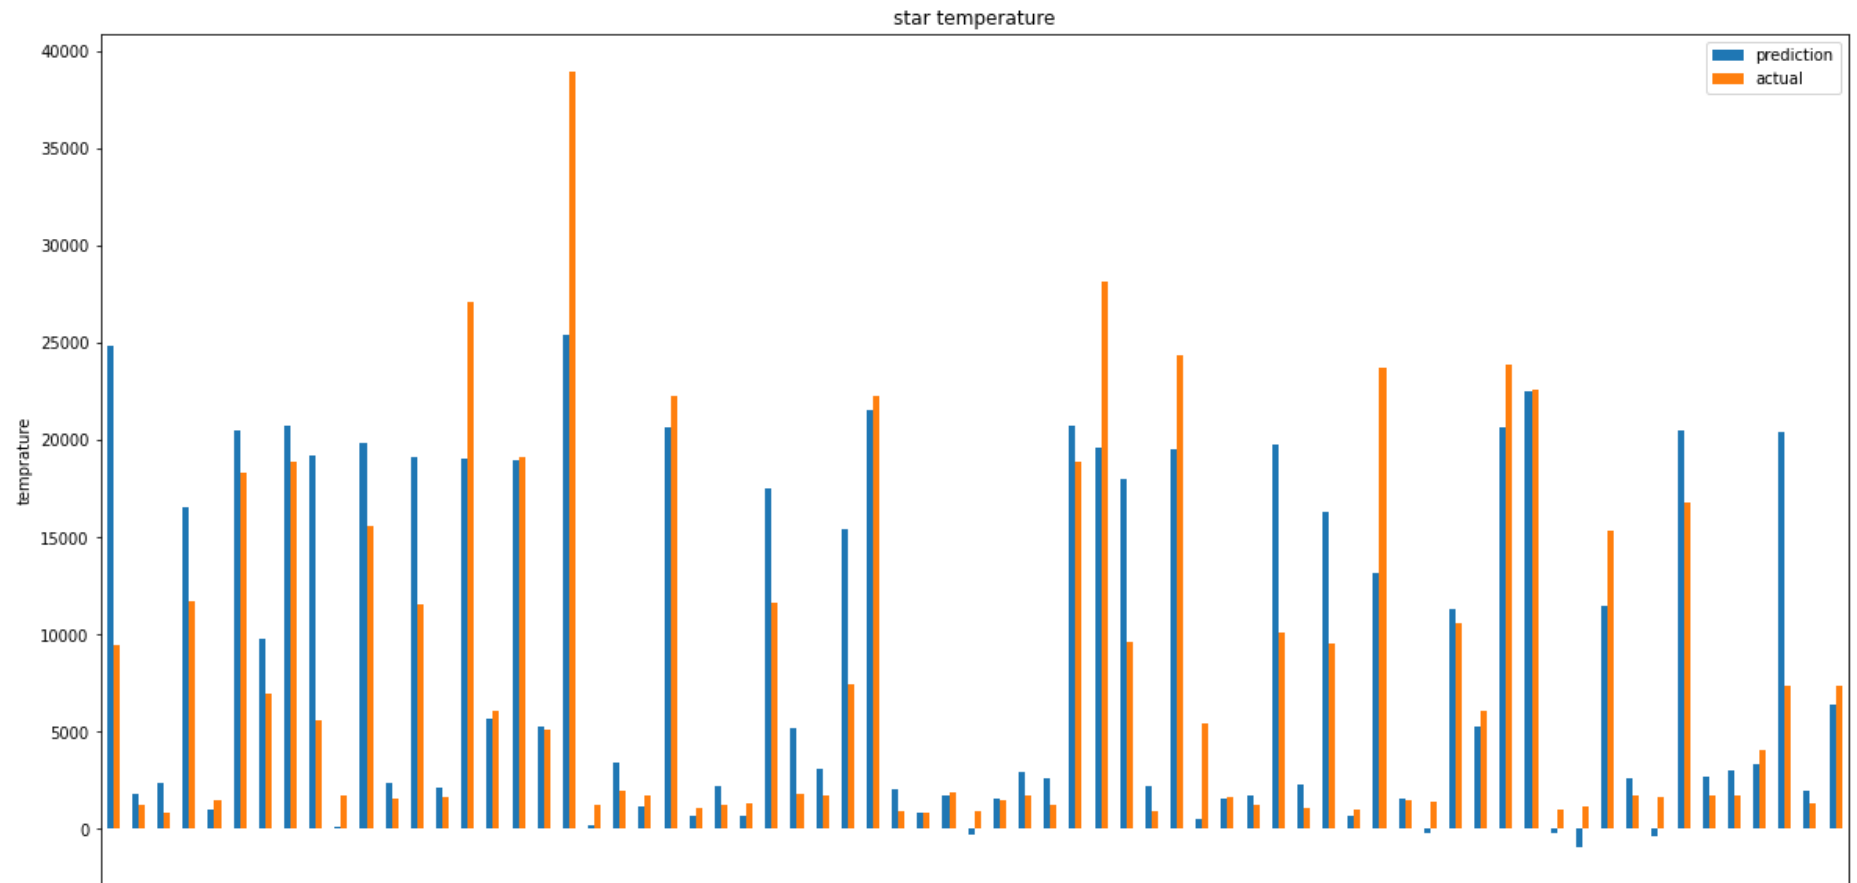
```

[illegible]

```
In [75]: plot_result(test_preds_4,y_test)
```


	prediction	actual
0	24865.683594	9449.041992
1	1804.027466	1272.693848
2	2340.869141	799.768799
3	16545.644531	11718.031250
4	977.693787	1472.373291

RMSE: 4803.796



Cocnlusion on dropout for net_4

Score increase obtained only in case of apply $dp1 = 0.2$

Other settings has not increase the RMS scoe for the net 4

In [76]: *# based on optimised net 2 building the net_5*

```
torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8
output_size = 1

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.dp1 = nn.Dropout(p = 0.1)
        self.act1 = nn.LeakyReLU()
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        # self.dp2 = nn.Dropout(p = 0.2)
        self.act2 = nn.LeakyReLU()
        self.fc3 = nn.Linear(hidden_size_2, output_size)
        # self.dp3 = nn.Dropout(p = 0.2)
        self.act3 = nn.LeakyReLU()

    def forward(self, x):
        x = self.fc1(x)
        x = self.dp1(x)
        x = self.act1(x)
        x = self.fc2(x)
        # x = self.dp2(x)
        x = self.act2(x)
        x = self.fc3(x)
        # x = self.dp3(x)
        x = self.act3(x)
        return x

net_5 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)
```

In [77]: *# training net_5 (changes - added dp1, p = 0.1)*

```
optimizer = torch.optim.Adam(net_5.parameters(), lr=0.001)

loss = torch.nn.MSELoss()

num_epochs = 500
```

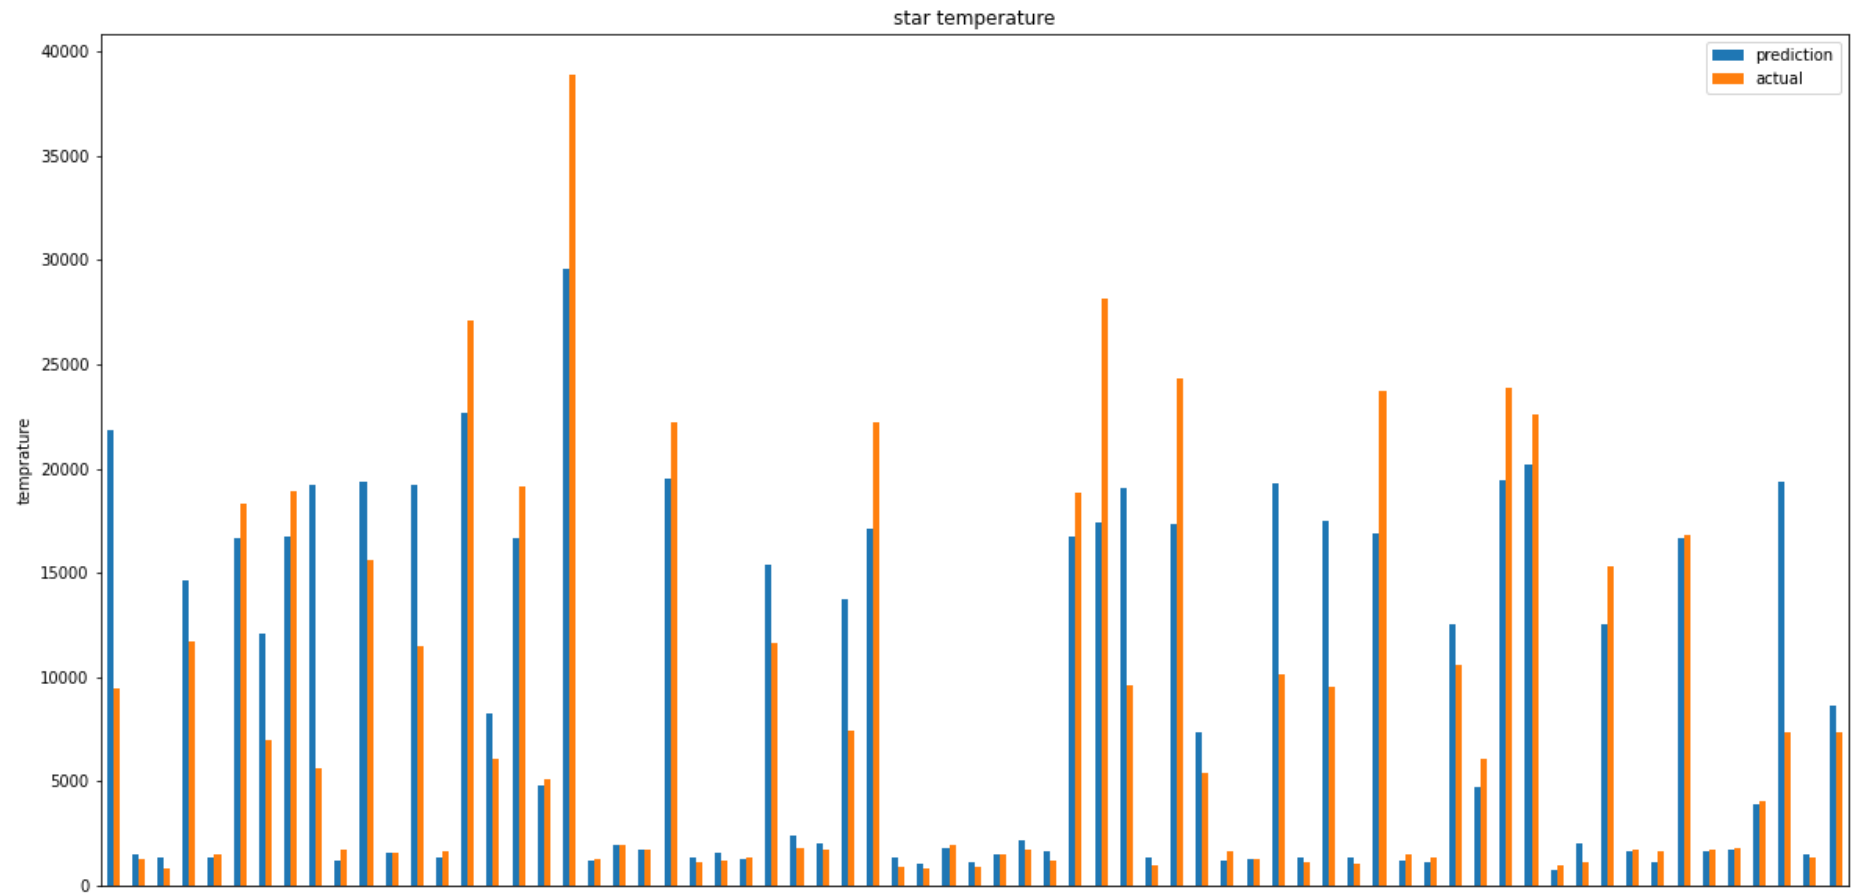
```
for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_5.forward(X_train).flatten()
    loss_value = loss(preds,y_train)
    loss_value.backward()
    optimizer.step()
    if (epoch % 25 == 0) or (epoch == 500) :
        print(loss_value)
        net_5.eval(),
        test_preds_5 = net_5.forward(X_test).flatten()
        accuracy = (torch.round(test_preds_5) == y_test).float().mean().data
        print(accuracy)
```

```
tensor(0.1119, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0810, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0614, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0445, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0313, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0251, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0220, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0197, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0179, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0163, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0149, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0141, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0135, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0131, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0128, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0126, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0124, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0122, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0119, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0117, grad_fn=<MseLossBackward0>)  
tensor(0.)
```

```
In [78]: plot_result(test_preds_5,y_test)
```

	prediction	actual
0	21826.535156	9449.041992
1	1492.371826	1272.693848
2	1367.966187	799.768799
3	14648.825195	11718.031250
4	1352.701904	1472.373291

RMSE: 4337.173



In [79]: *# based on optimised net 2 building the net_5*

```

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8

```

```

output_size = 1

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.dp1 = nn.Dropout(p = 0.1)
        self.act1 = nn.LeakyReLU()
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.dp2 = nn.Dropout(p = 0.1)
        self.act2 = nn.LeakyReLU()
        self.fc3 = nn.Linear(hidden_size_2, output_size)
        # self.dp3 = nn.Dropout(p = 0.2)
        self.act3 = nn.LeakyReLU()

    def forward(self, x):
        x = self.fc1(x)
        x = self.dp1(x)
        x = self.act1(x)
        x = self.fc2(x)
        x = self.dp2(x)
        x = self.act2(x)
        x = self.fc3(x)
        # x = self.dp3(x)
        x = self.act3(x)
        return x

net_5 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)

```

In [80]: *# training net_5 (changes - added dp1, p = 0.1 , dp2, p=0.1)*

```

optimizer = torch.optim.Adam(net_5.parameters(),lr=0.001)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_5.forward(X_train).flatten()
    loss_value = loss(preds,y_train)
    loss_value.backward()
    optimizer.step()

```

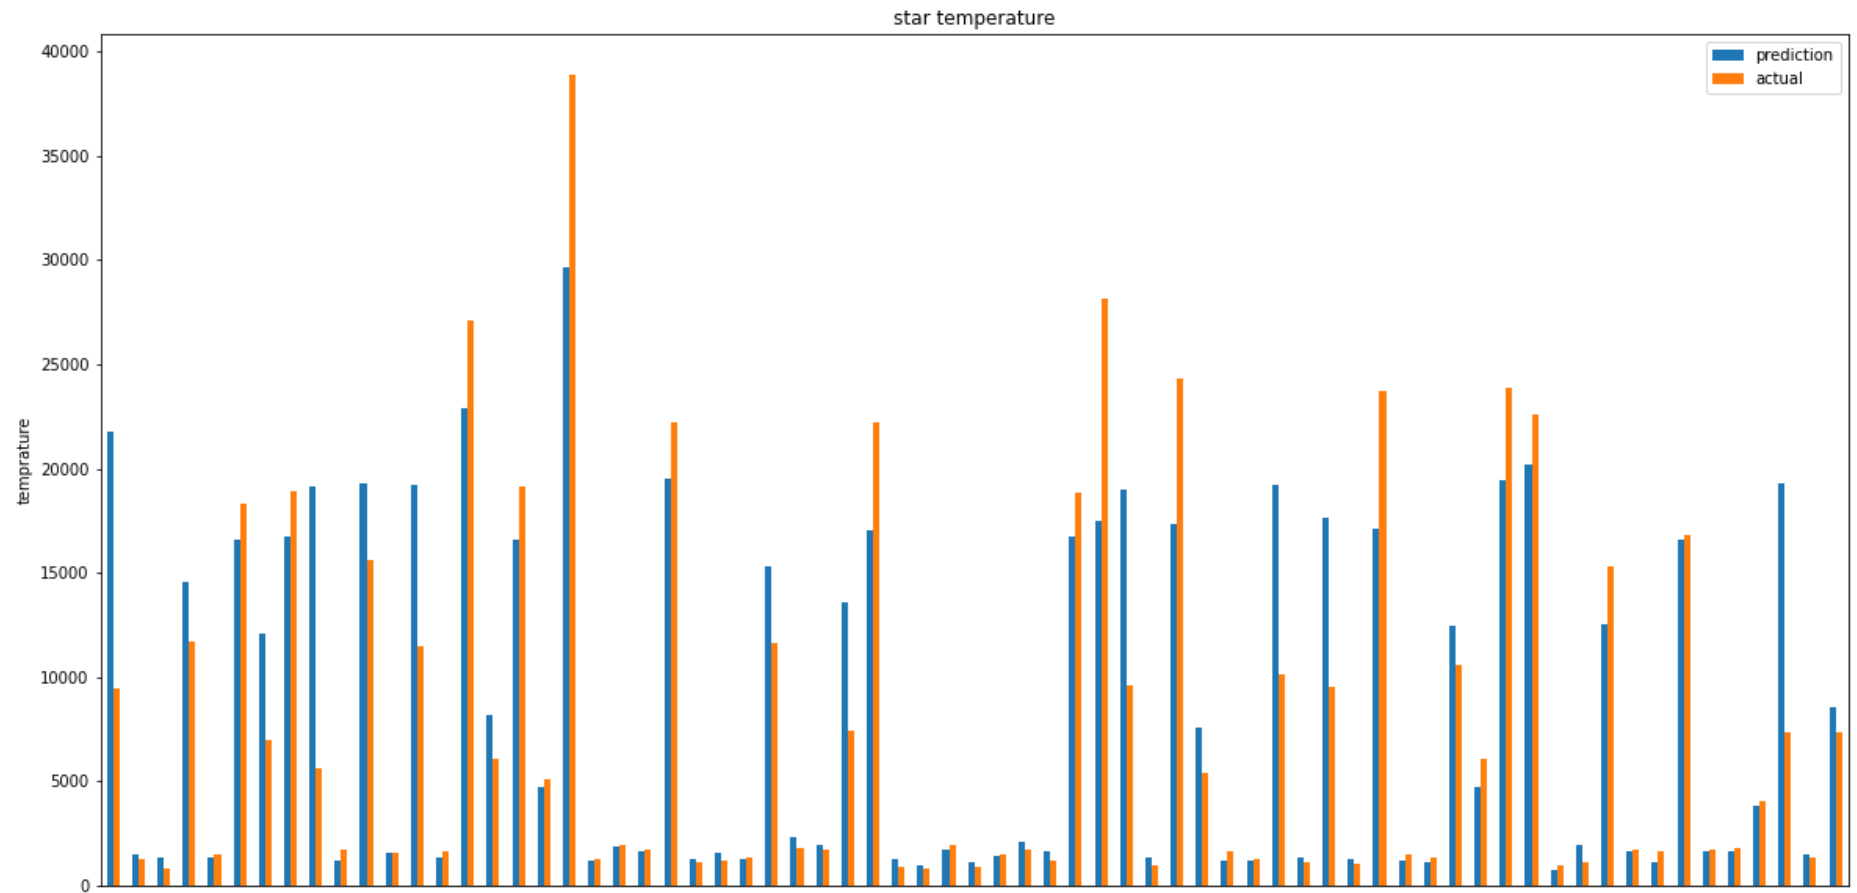
```
if (epoch % 25 == 0) or (epoch == 500) :  
    print(loss_value)  
    net_5.eval(),  
    test_preds_5 = net_5.forward(X_test).flatten()  
    accuracy = (torch.round(test_preds_5) == y_test).float().mean().data  
    print(accuracy)
```

```
tensor(0.1130, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0809, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0614, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0444, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0313, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0250, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0219, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0196, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0178, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0162, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0149, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0140, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0135, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0131, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0128, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0126, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0124, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0121, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0119, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0117, grad_fn=<MseLossBackward0>)  
tensor(0.)
```

```
In [81]: plot_result(test_preds_5,y_test)
```


	prediction	actual
0	21789.283203	9449.041992
1	1470.095215	1272.693848
2	1351.251099	799.768799
3	14539.531250	11718.031250
4	1329.332153	1472.373291

RMSE: 4318.3047



In [82]: *# based on optimised net 2 building the net_5*

```

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8

```

```

output_size = 1

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.dp1 = nn.Dropout(p = 0.1)
        self.act1 = nn.LeakyReLU()
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.dp2 = nn.Dropout(p = 0.05)
        self.act2 = nn.LeakyReLU()
        self.fc3 = nn.Linear(hidden_size_2, output_size)
        self.dp3 = nn.Dropout(p = 0.05)
        self.act3 = nn.LeakyReLU()

    def forward(self, x):
        x = self.fc1(x)
        x = self.dp1(x)
        x = self.act1(x)
        x = self.fc2(x)
        x = self.dp2(x)
        x = self.act2(x)
        x = self.fc3(x)
        x = self.dp3(x)
        x = self.act3(x)
        return x

net_5 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)

```

```

In [83]: # training net_5 (changes - added dp1, p = 0.1 , dp2, p=0.05, dp3 = 0.05)

optimizer = torch.optim.Adam(net_5.parameters(),lr=0.001)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_5.forward(X_train).flatten()
    loss_value = loss(preds,y_train)
    loss_value.backward()
    optimizer.step()

```

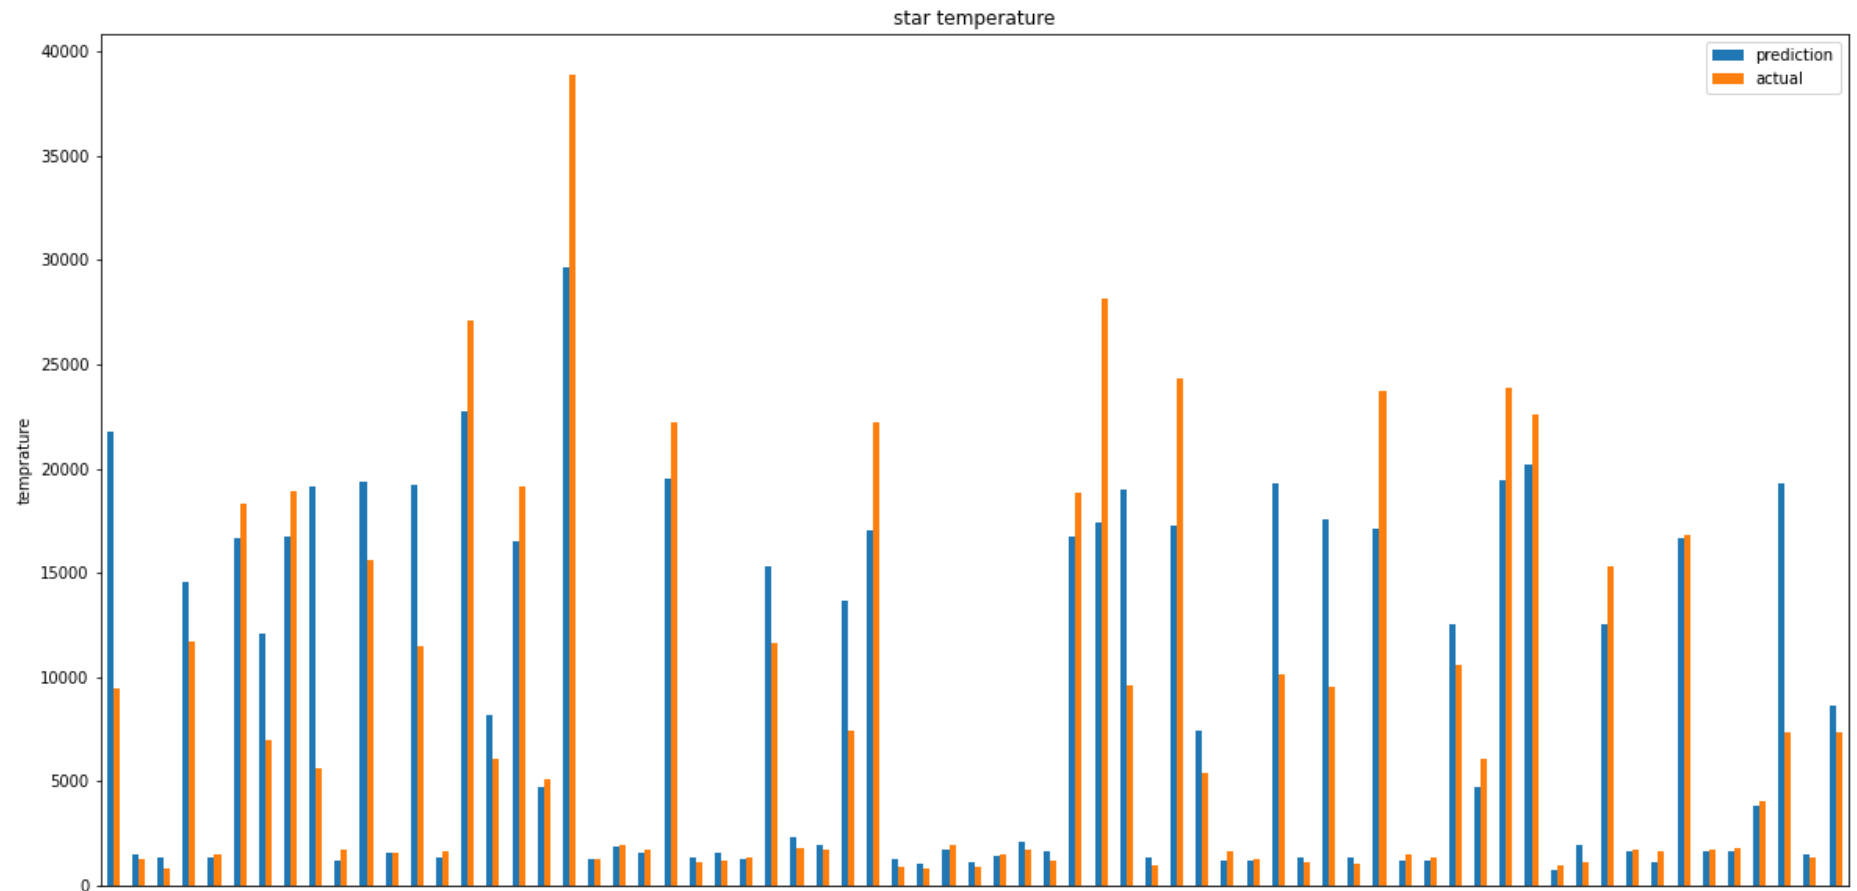
```
if (epoch % 25 == 0) or (epoch == 500) :  
    print(loss_value)  
    net_5.eval(),  
    test_preds_5 = net_5.forward(X_test).flatten()  
    accuracy = (torch.round(test_preds_5) == y_test).float().mean().data  
    print(accuracy)
```

```
tensor(0.1126, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0809, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0614, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0445, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0313, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0251, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0219, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0196, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0178, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0162, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0148, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0140, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0135, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0131, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0128, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0126, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0124, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0122, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0119, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0117, grad_fn=<MseLossBackward0>)  
tensor(0.)
```

```
In [84]: plot_result(test_preds_5,y_test)
```

	prediction	actual
0	21795.246094	9449.041992
1	1478.717285	1272.693848
2	1351.359741	799.768799
3	14566.385742	11718.031250
4	1346.858398	1472.373291

RMSE: 4328.24



Conclusion on dropout tuning on net_5

Best score achieved adding the following layers:

- dp1 p = 0.1
- dp2 p = 0.05

- $dp3\ p = 0.05$

The changes to be used for further network optimisation

BatchNorm Tuning

```
In [85]: # based on optimised net 1 building the net_4

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8
output_size = 1

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.dp1 = nn.Dropout(p = 0.2)
        self.bn1 = nn.BatchNorm1d(hidden_size_1)
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.fc3 = nn.Linear(hidden_size_2, output_size)

    def forward(self, x):
        x = self.fc1(x)
        x = self.dp1(x)
        x = self.bn1(x)
        x = self.fc2(x)
        x = self.fc3(x)
        return x

net_4 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)
```

```
In [86]: # training net_4 (changes - added bn1)

optimizer = torch.optim.Adam(net_4.parameters(), lr=0.011)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
```


```
optimizer.zero_grad()
preds = net_4.forward(X_train).flatten()
loss_value = loss(preds,y_train)
loss_value.backward()
optimizer.step()
if (epoch % 25 == 0) or (epoch == 500) :
    print(loss_value)
    net_4.eval(),
    test_preds_4 = net_4.forward(X_test).flatten()
    accuracy = (torch.round(test_preds_4) == y_test).float().mean().data
    print(accuracy)
```

```
tensor(0.3048, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0194, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0164, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0161, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0158, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0158, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0158, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0158, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0158, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0157, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0157, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0157, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0157, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0156, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0156, grad_fn=<MseLossBackward0>)  
tensor(0.)
```

```
In [87]: plot_result(test_preds_4,y_test)
```


	prediction	actual
0	25039.654297	9449.041992
1	1715.798096	1272.693848
2	2056.192139	799.768799
3	16910.980469	11718.031250
4	1071.918335	1472.373291

RMSE: 4807.141



In [88]: *# based on optimised net 1 building the net_4*

```

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8
output_size = 1

```

```

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.dp1 = nn.Dropout(p = 0.2)
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.bn2 = nn.BatchNorm1d(hidden_size_2)
        self.fc3 = nn.Linear(hidden_size_2, output_size)

    def forward(self, x):
        x = self.fc1(x)
        x = self.dp1(x)
        x = self.fc2(x)
        x = self.bn2(x)
        x = self.fc3(x)
        return x

net_4 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)

```

In [89]: *# training net_4 (changes - added bn2)*

```

optimizer = torch.optim.Adam(net_4.parameters(), lr=0.011)

loss = torch.nn.MSELoss()

num_epochs = 500

for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_4.forward(X_train).flatten()
    loss_value = loss(preds, y_train)
    loss_value.backward()
    optimizer.step()
    if (epoch % 25 == 0) or (epoch == 500) :
        print(loss_value)
        net_4.eval(),
        test_preds_4 = net_4.forward(X_test).flatten()
        accuracy = (torch.round(test_preds_4) == y_test).float().mean().data
        print(accuracy)

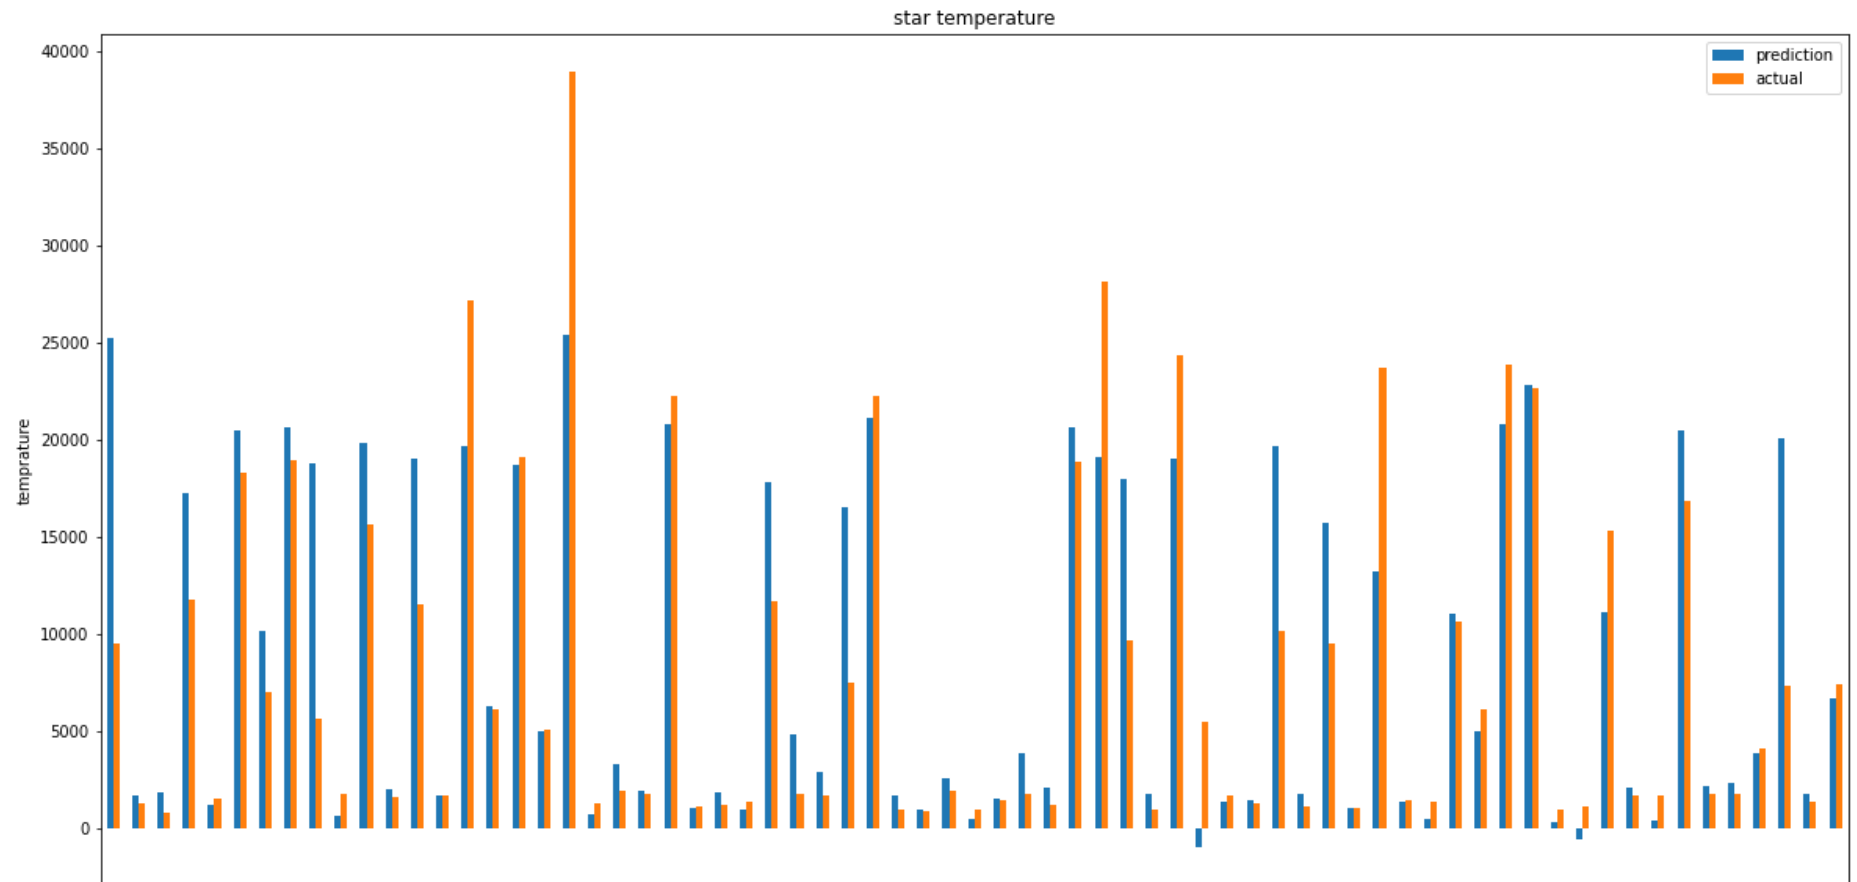
```

```
tensor(0.5774, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0244, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0165, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0161, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0161, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0160, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0159, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0158, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0158, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0158, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0158, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0158, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0158, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0157, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0157, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0157, grad_fn=<MseLossBackward0>)  
tensor(0.)
```

```
In [90]: plot_result(test_preds_4,y_test)
```

	prediction	actual
0	25218.388672	9449.041992
1	1640.124023	1272.693848
2	1796.611572	799.768799
3	17222.617188	11718.031250
4	1156.061279	1472.373291

RMSE: 4819.2676



Conclusion on using of Batchnorm on net_4

Best score network has achived with applied bn2.

In [91]: *# based on optimised net 2 building the net_5*

```

torch.manual_seed(1234)
input_size = 14
hidden_size_1 = 12
hidden_size_2 = 8
output_size = 1

class NeuralNet(nn.Module):
    def __init__(self, input_size, hidden_size_1, hidden_size_2, output_size):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size_1)
        self.dp1 = nn.Dropout(p = 0.1)
        self.act1 = nn.LeakyReLU()
        self.bn1 = nn.BatchNorm1d(hidden_size_1)
        self.fc2 = nn.Linear(hidden_size_1, hidden_size_2)
        self.dp2 = nn.Dropout(p = 0.05)
        self.act2 = nn.LeakyReLU()
        self.bn2 = nn.BatchNorm1d(hidden_size_2)
        self.fc3 = nn.Linear(hidden_size_2, output_size)
        self.dp3 = nn.Dropout(p = 0.05)
        self.act3 = nn.LeakyReLU()

    def forward(self, x):
        x = self.fc1(x)
        x = self.dp1(x)
        x = self.act1(x)
        x = self.bn1(x)
        x = self.fc2(x)
        x = self.dp2(x)
        x = self.act2(x)
        x = self.bn2(x)
        x = self.fc3(x)
        x = self.dp3(x)
        x = self.act3(x)
        return x

net_5 = NeuralNet(input_size, hidden_size_1, hidden_size_2, output_size)

```

```

In [92]: # training net_5 (changes - added bn1 u bn2)

optimizer = torch.optim.RMSprop(net_5.parameters(), lr=0.001)

loss = torch.nn.MSELoss()

```

```
num_epochs = 500


for epoch in range(num_epochs):
    optimizer.zero_grad()
    preds = net_5.forward(X_train).flatten()
    loss_value = loss(preds, y_train)
    loss_value.backward()
    optimizer.step()
    if (epoch % 25 == 0) or (epoch == 500) :
        print(loss_value)
        net_5.eval(),
        test_preds_5 = net_5.forward(X_test).flatten()
        accuracy = (torch.round(test_preds_5) == y_test).float().mean().data
        print(accuracy)
```

```
tensor(0.1984, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0548, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0319, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0224, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0192, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0170, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0155, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0145, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0139, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0135, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0132, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0129, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0127, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0126, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0124, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0124, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0123, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0121, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0121, grad_fn=<MseLossBackward0>)  
tensor(0.)  
tensor(0.0120, grad_fn=<MseLossBackward0>)  
tensor(0.)
```

```
In [93]: plot_result(test_preds_5,y_test)
```

	prediction	actual
0	21879.076172	9449.041992
1	1796.710449	1272.693848
2	1542.364990	799.768799
3	15064.061523	11718.031250
4	1692.044678	1472.373291

RMSE: 4300.2734



Conclusion on application of Batchnorm on network_5

Using Batchnorm on first and secon layers model has slightly increase score.

General Conclusion

Creation of table with neural networks and it's scores

```
In [94]: table = []  
for i in range(5):  
    table.append('net_' + str(i+1))  
  
print(table)
```

```
['net_1', 'net_2', 'net_3', 'net_4', 'net_5']
```

```
In [95]: preds_nets = [test_preds_1, test_preds_2, test_preds_3, test_preds_4, test_preds_5]
```

```
In [96]: for i in range(len(preds_nets)):  
    preds_nets[i] = preds_nets[i].detach().numpy()
```

```
In [97]: rmse_nets = []  
for i in preds_nets:  
    rmse_nets.append(mean_squared_error(y_test.detach().numpy() * df_stars['Temperature (K)'].max(), i, squared = False))
```

```
In [98]: df_result = pd.DataFrame(table, columns = ['net_name'])
```

```
In [99]: df_result['prediction'] = preds_nets
```

```
In [100]: df_result['rmse'] = rmse_nets
```

During the project execution the data was imported, prepared and analyzed.

Five neural networks wer builded with following scores:

```
In [101]: df_result.sort_values(by = 'rmse', ascending = False)
```

Out[101]:

	net_name	prediction	rmse
0	net_1	[25638.512, 1454.7306, 1245.5305, 17902.568, 1...	4868.386230
1	net_2	[21669.697, 1492.6152, 1537.0199, 14472.895, 1...	4356.758301
2	net_3	[21032.305, 21063.041, 21090.668, 21094.648, 2...	15818.150391
3	net_4	[25218.389, 1640.124, 1796.6116, 17222.617, 11...	4819.267578
4	net_5	[21879.076, 1796.7104, 1542.365, 15064.062, 16...	4300.273438

Best score of neural network:

In [102... `df_result[df_result['rmse'] == df_result['rmse'].min()]`

Out[102]:

	net_name	prediction	rmse
4	net_5	[21879.076, 1796.7104, 1542.365, 15064.062, 16...	4300.273438

Best scores was achivied in neural network No 5 with following parameters:

- lr =0.01
- num_epochs = 500
- optimizer RMS PROP
- activation function - nn.LeakyReLU()
- Dropout(p = 0.1) on first layer and p = 0.05 second and third layers;
- batchnorm - on first and second layers.

RMSE less than 4500 successfully achieved!